



# Python for Data Science using Anaconda

TDWI Accelerate

Travis E. Oliphant  
*President, Chief Data Scientist, Co-founder*

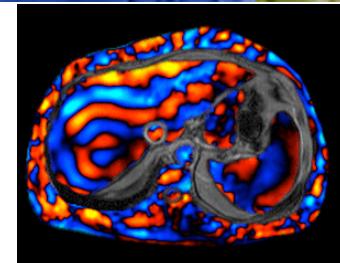
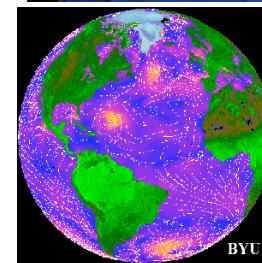
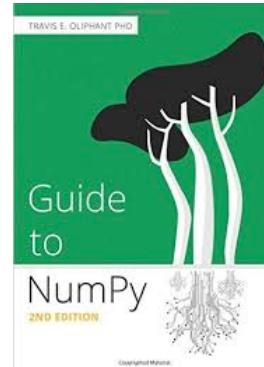


© 2016 Continuum Analytics - Confidential & Proprietary



# A bit about me

- PhD 2001 from Mayo Clinic in Biomedical Engineering (Ultrasound and MRI)
- MS/BS degrees in Elec. Comp. Engineering
- Creator and Developer of **SciPy** (1999-2009)
- Professor at BYU (2001-2007)
- Author of **NumPy** (2005-2012)
- Started **Numba** (2012)
- Founding Chair of **NumFOCUS / PyData**
- Previous Python Software Foundation Director
- Co-founder of Continuum Analytics
- CEO => President, Chief Data Scientist



SciPy



# 2012 - Created Two Orgs for Sustainable Open Source

## Company

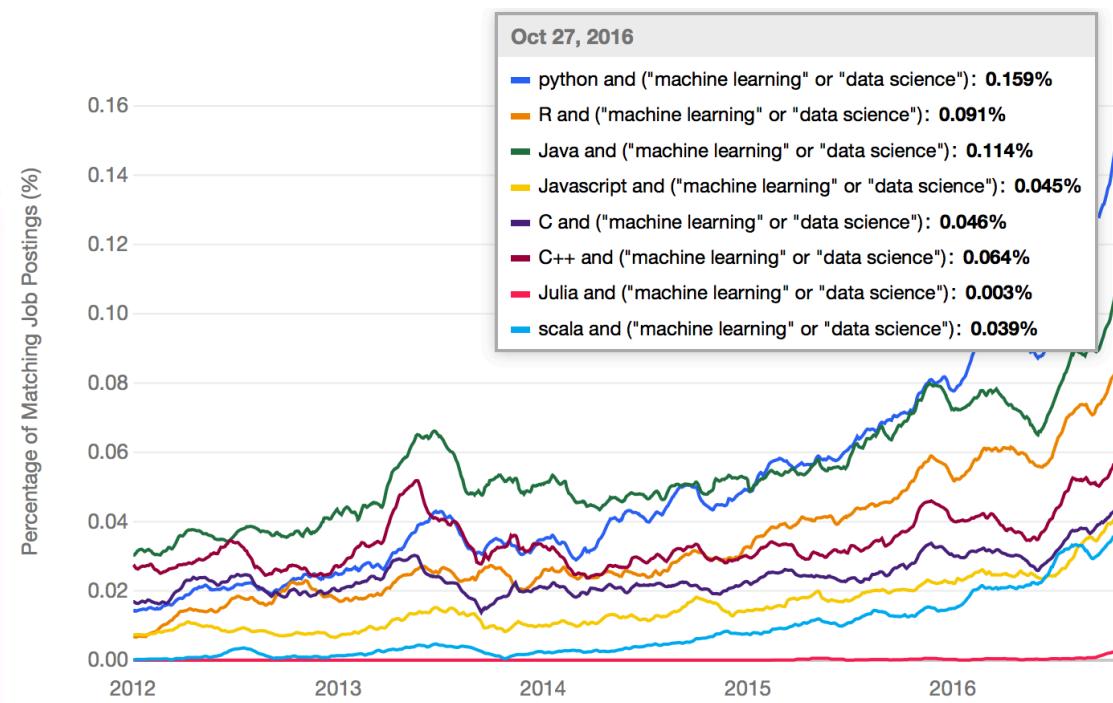
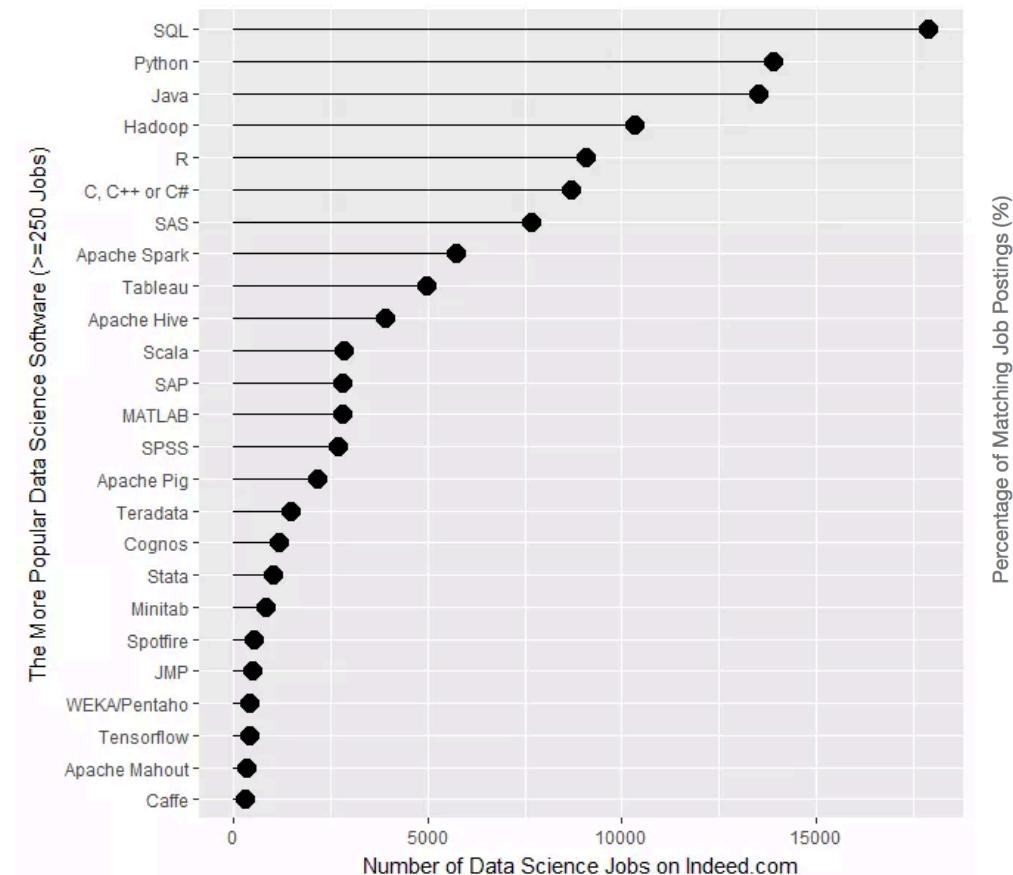


Enterprise **software, support and services** to empower people who change the world to get rapid insight from **all** of their data — built on **open-source** that we contribute to and sustain.

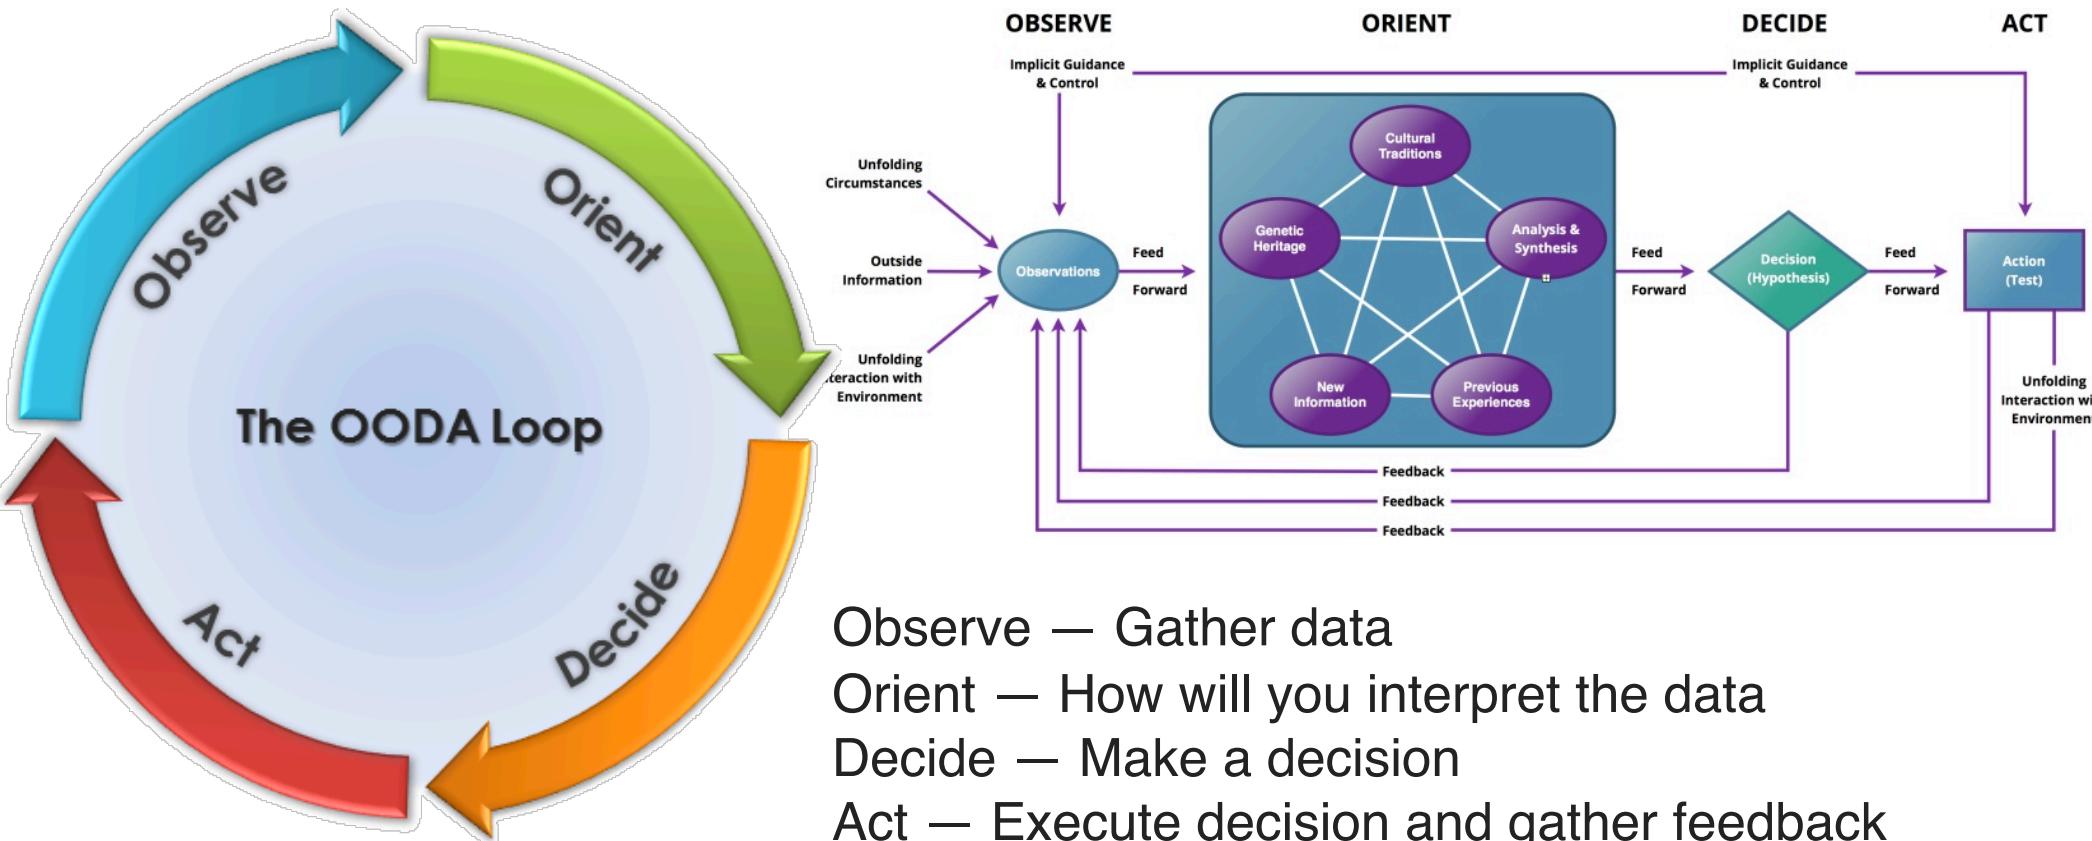
## Community



# Data Science is growing with Python leading the way

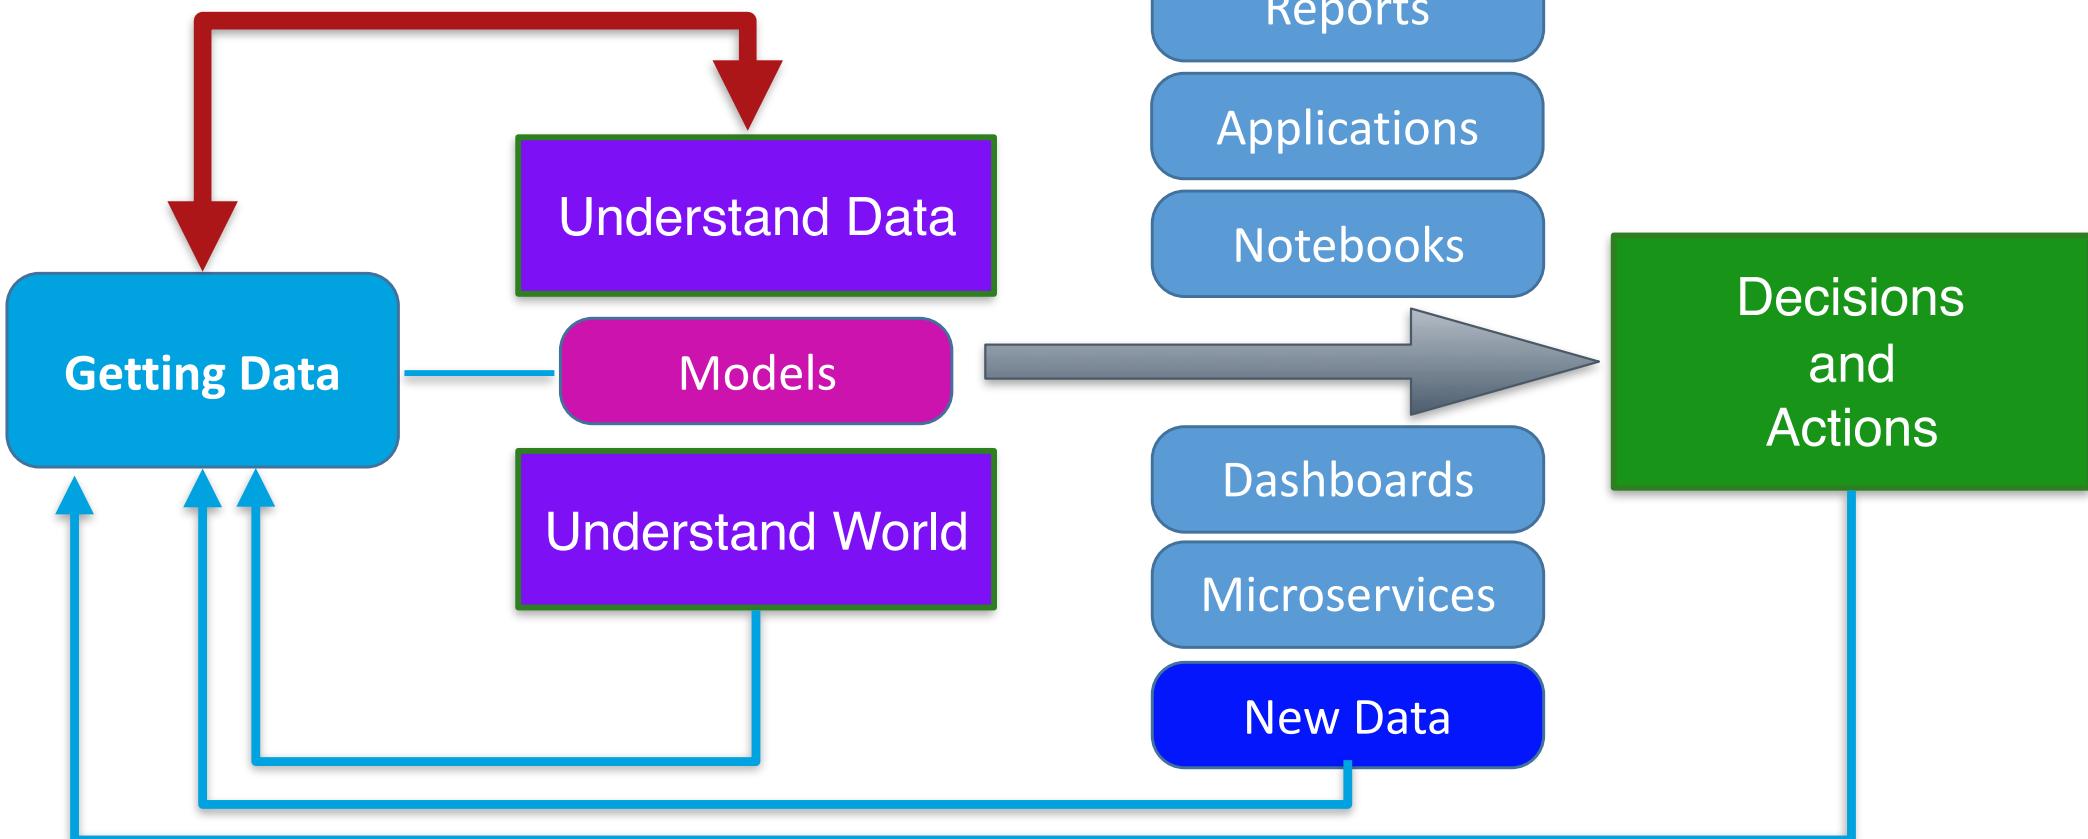


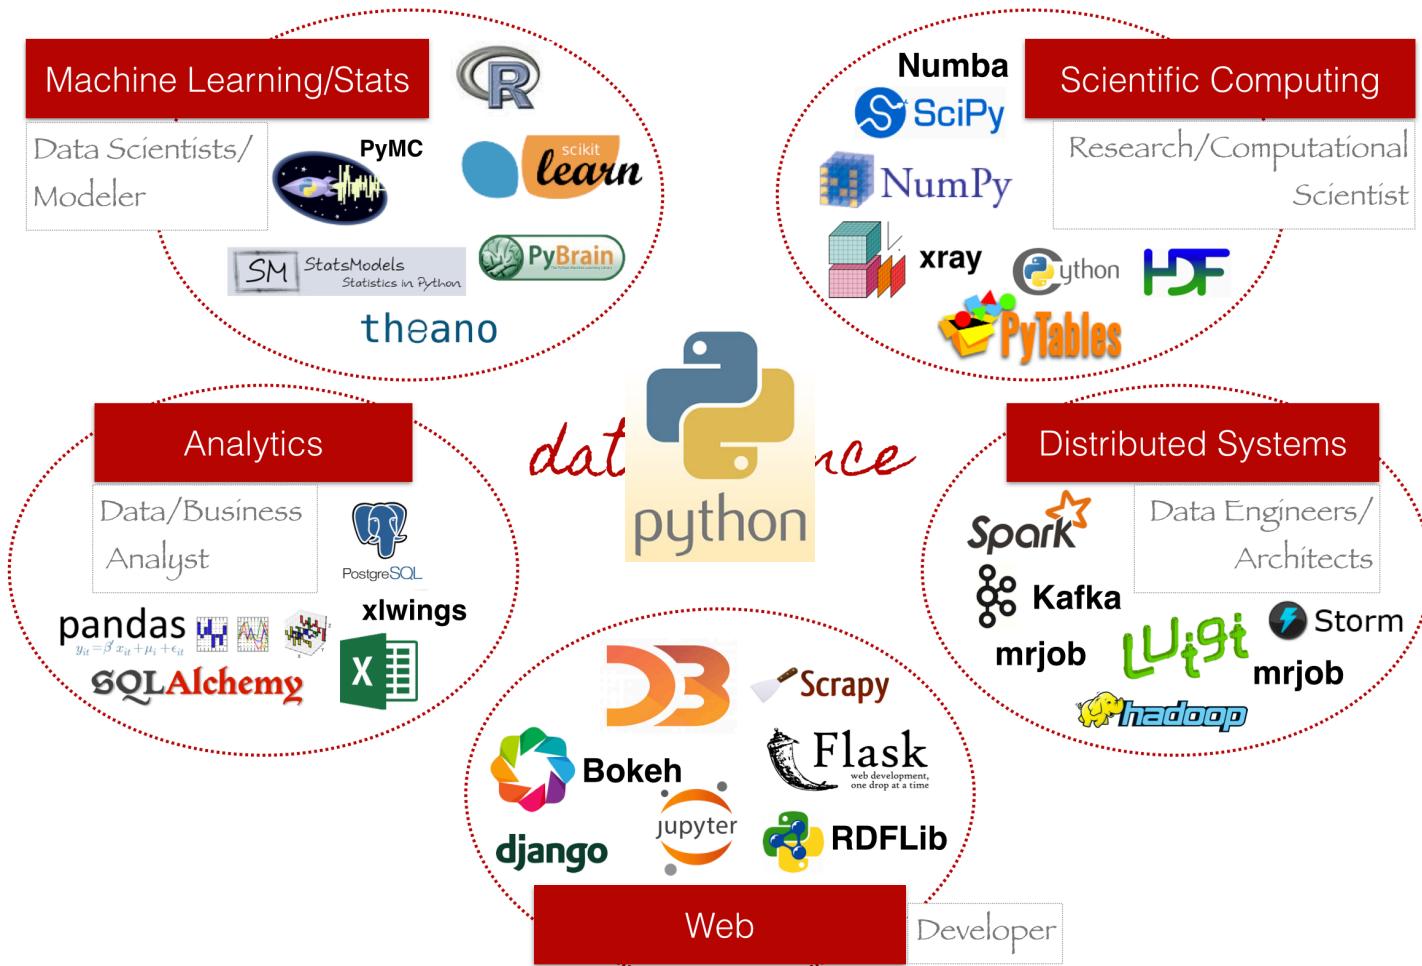
# Dealing with a lot of information: OODA loop



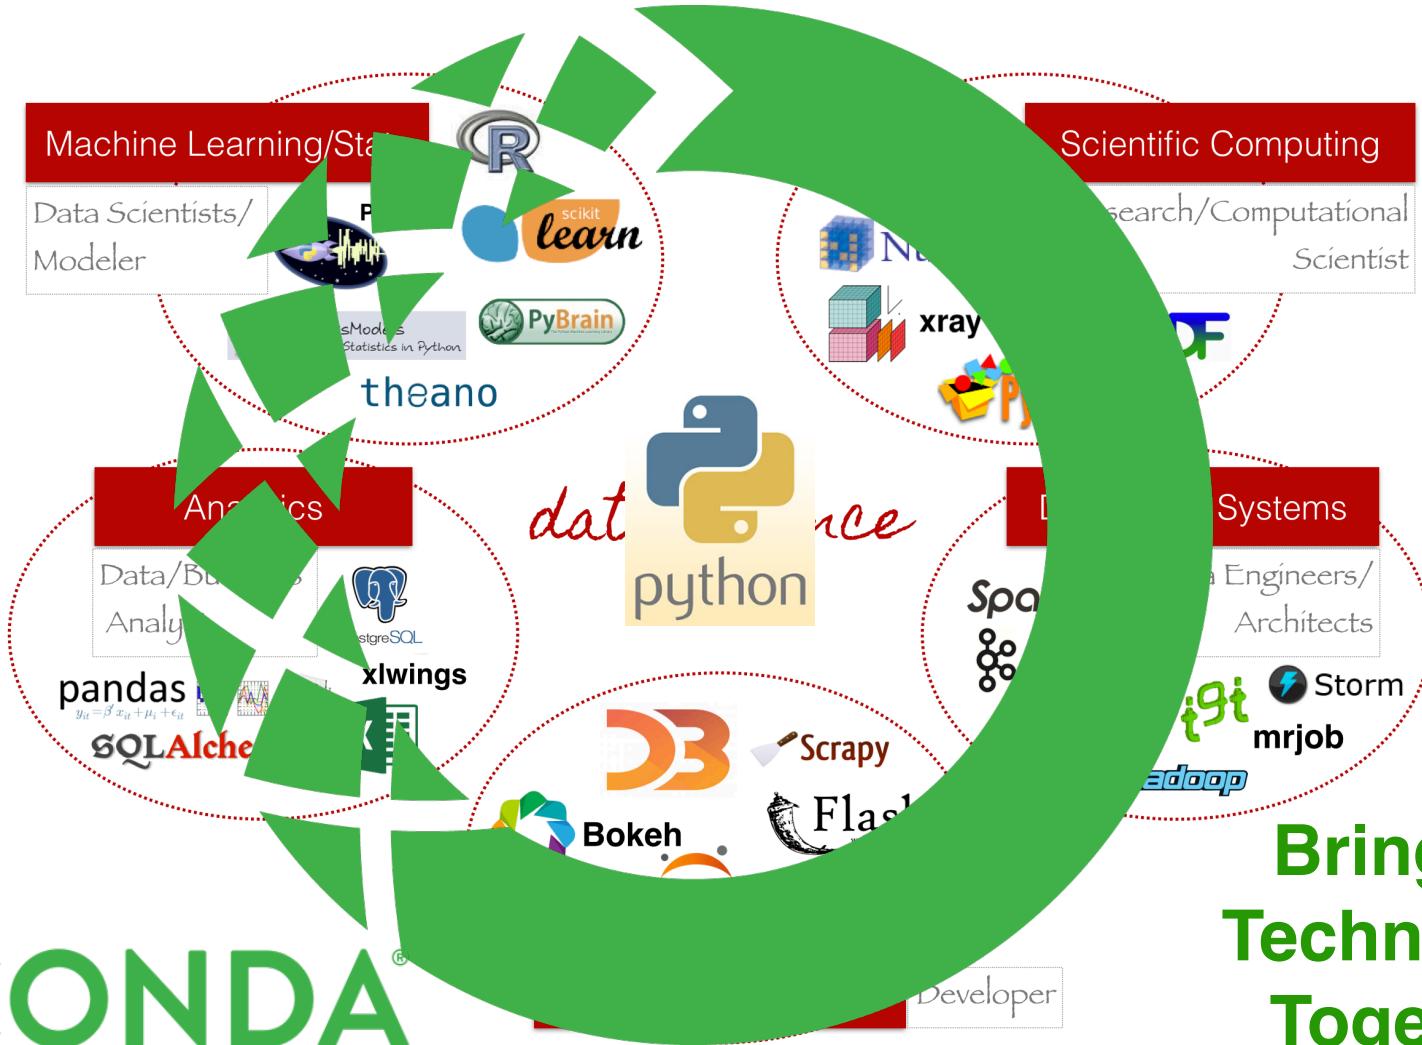
# Data Science Workflow

## Exploratory Data Analysis and Viz





# ANACONDA



**Bringing  
Technology  
Together**

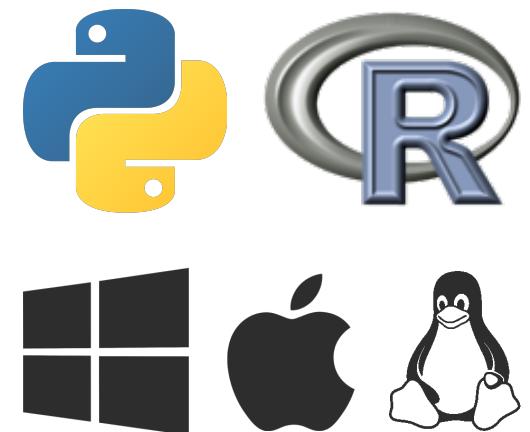
# Accelerate adoption of Data Science

<http://continuum.io/downloads>



PYTHON & R OPEN SOURCE ANALYTICS

NumPy	SciPy	Pandas	Scikit-learn	Jupyter/IPython
Numba	Matplotlib	Spyder	TensorFlow	Cython    Theano
Scikit-image	NLTK	Dask	Caffe	dplyr    shiny
ggplot2	tidyr	caret	PySpark	& 1000+ packages
<b>CONDA®</b>				



# ANACONDA

## Trusted by Industry Leaders

### Financial Services

Risk Mgmt, Quant modeling, Data exploration and processing, algorithmic trading, compliance reporting

### Government

Fraud detection, data crawling, web & cyber data analytics, statistical modeling

### Healthcare & Life Sciences

Genomics data processing, cancer research, natural language processing for health data science

### High Tech

Customer behavior, recommendations, ad bidding, retargeting, social media analytics

### Retail & CPG

Engineering simulation, supply chain modeling, scientific analysis

### Oil & Gas

Pipeline monitoring, noise logging, seismic data processing, geophysics

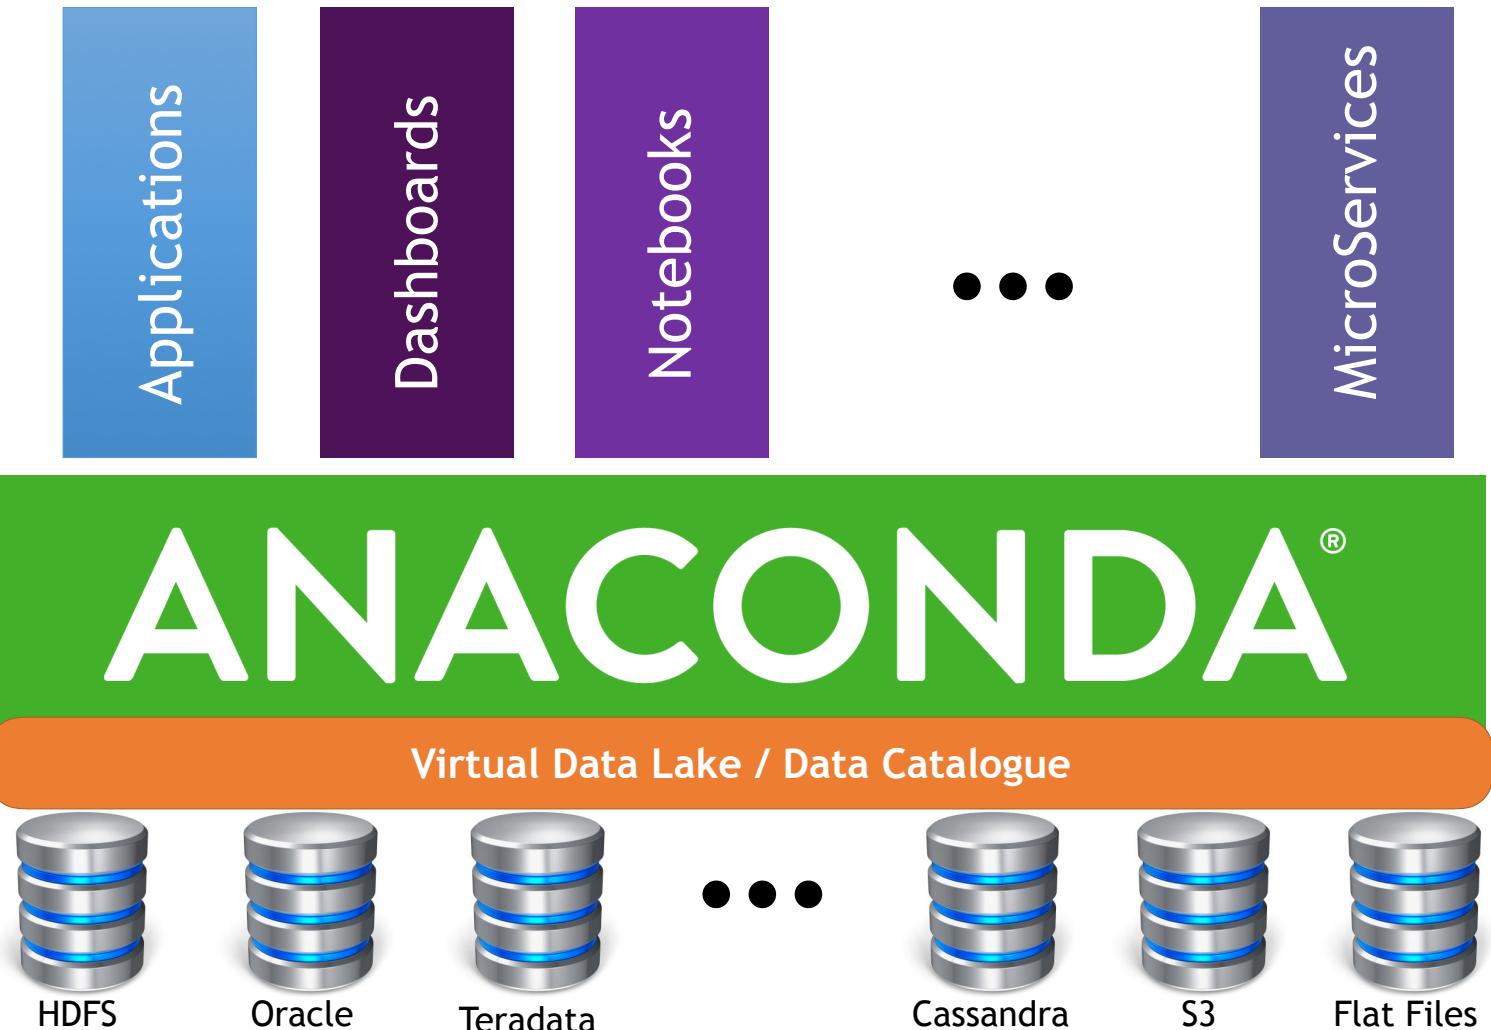
<http://continuum.io/downloads>

**Insight  
Decisions  
Actions  
Results  
Creations**

**Open Data  
Science**

**DATA**





# ANACONDA NAVIGATOR

[Home](#)[Environments](#)[Learning](#)[Community](#)[Documentation](#)[Developer Blog](#)[Feedback](#)

Applications on

rbokeh

Channels



notebook

4.2.3

Web-based, interactive computing notebook environment. Edit and run human-readable docs while describing the data analysis.

[Launch](#)

qtconsole

4.2.1

PyQt GUI that supports inline figures, proper multiline editing with syntax highlighting, graphical calltips, and more.

[Launch](#)

spyder

3.0.0

Scientific PYthon Development EnviRonment. Powerful Python IDE with advanced editing, interactive testing, debugging and introspection features

[Launch](#)

anaconda-fusion

1.0.2



glueviz

0.9.1



rstudio

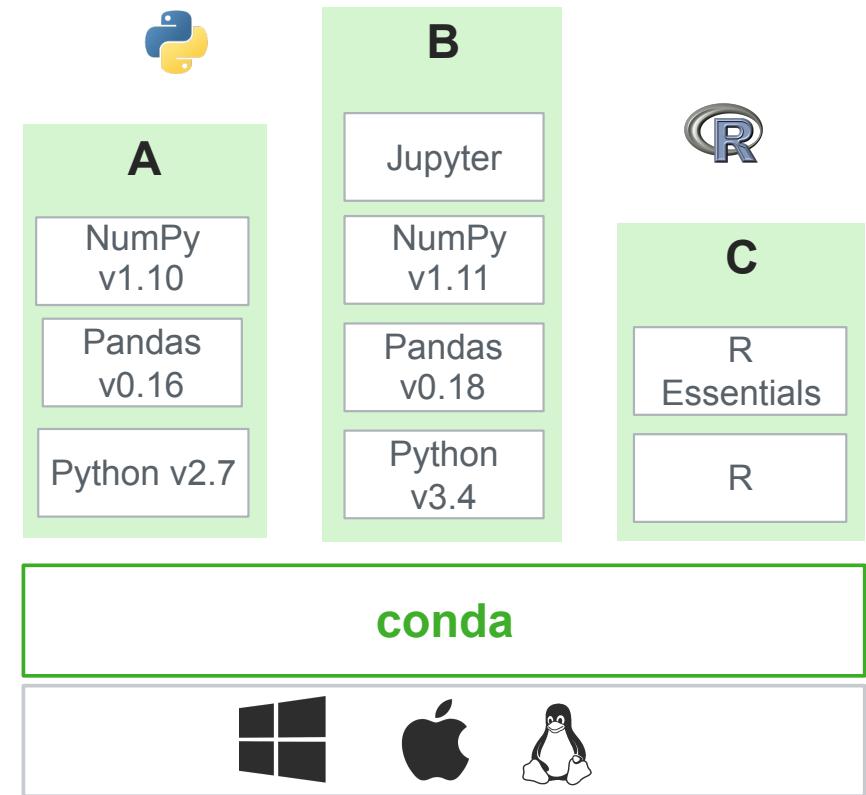
1.0.44



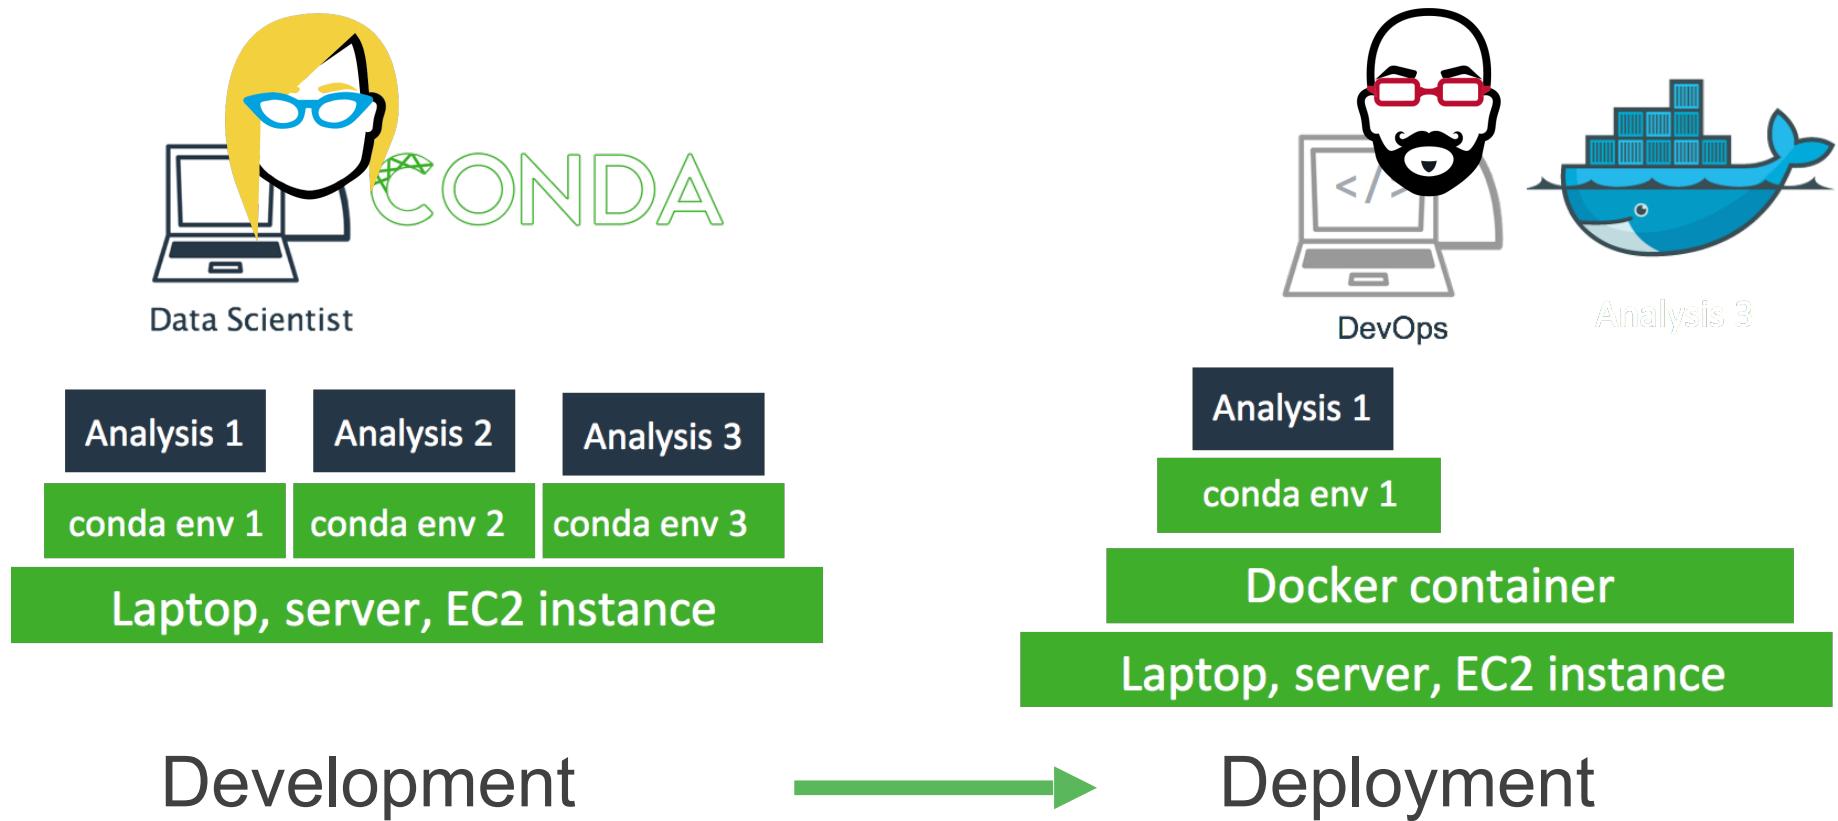
# Conda Sandboxing Technology

- **Language** independent
- **Platform** independent
- No special privileges required
- No VMs or containers
- Enables:
  - **Reproducibility**
  - **Collaboration**
  - **Scaling**

***“conda – package everything”***



# Conda for Deployment



# CONDA

```
$ conda install python=2.7  
$ conda install pandas  
$ conda install -c r r  
$ conda install -c conda-forge tensorflow
```

Install dependencies

```
$ anaconda-project run plot --show
```

Deploy an interactive visualization



© 2016 Continuum Analytics - Confidential & Proprietary

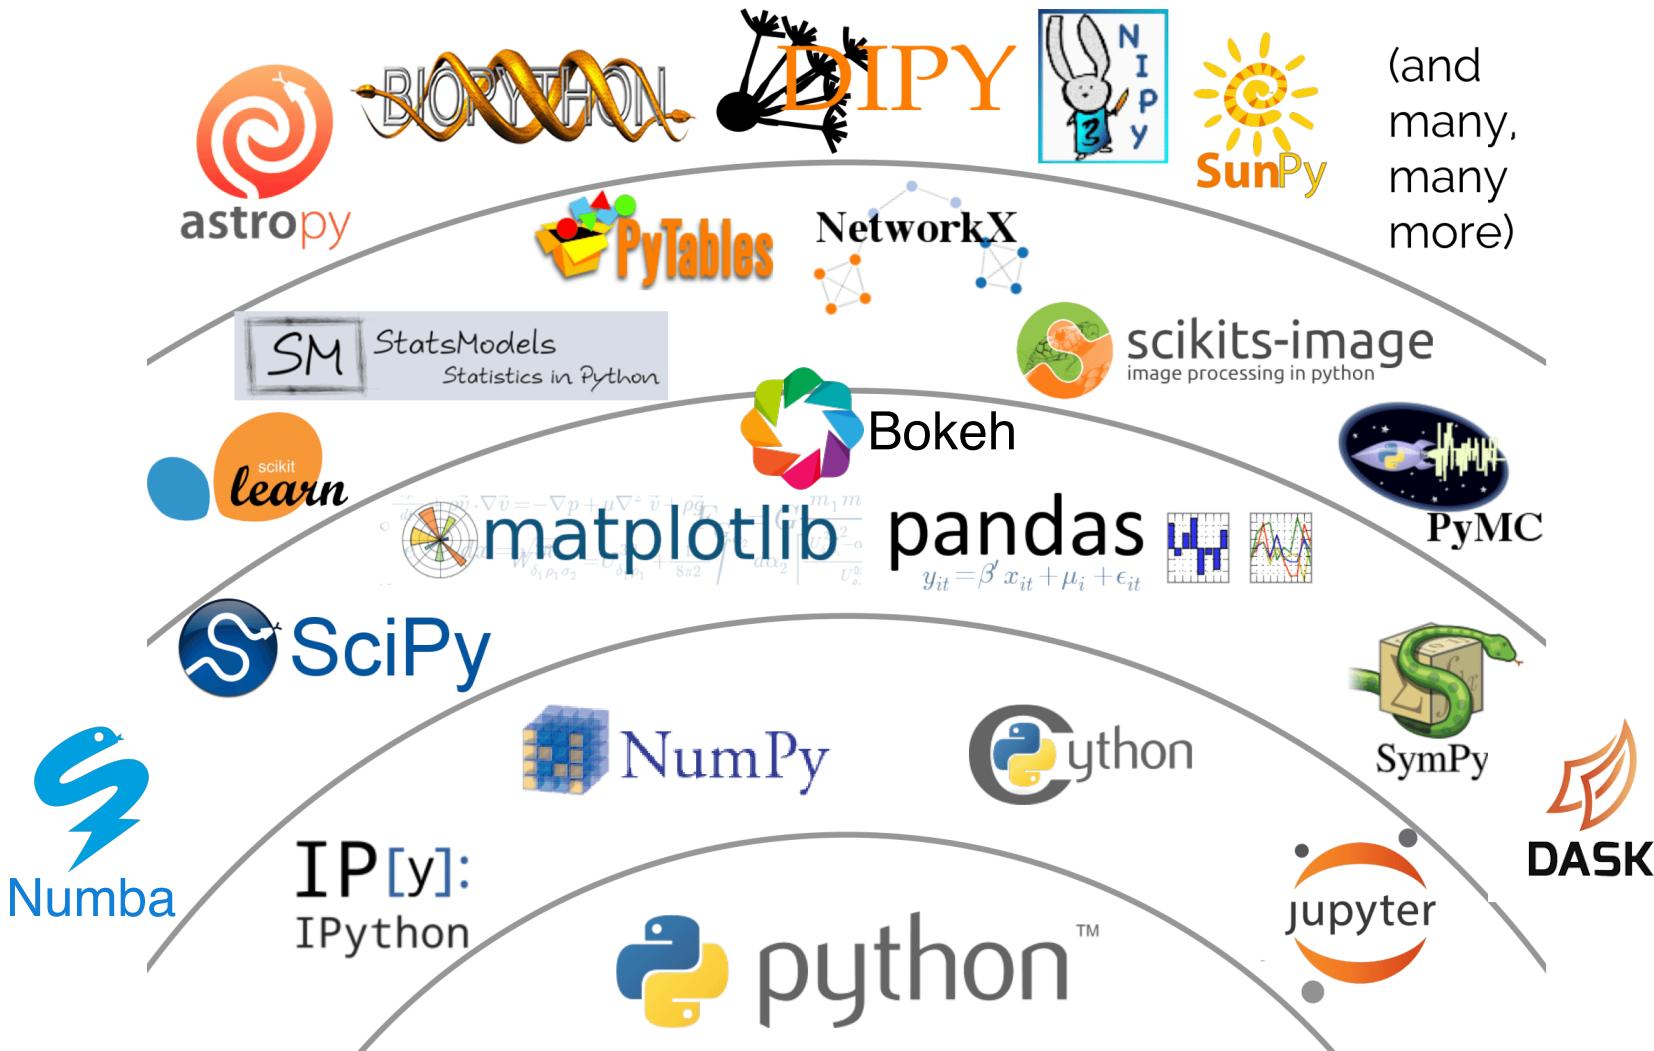
```
name: myenv  
channels:  
  - chdoig  
  - r  
  - foo  
dependencies:  
  - python=2.7  
  - r  
  - r-ldavis  
  - pandas  
  - mongodb  
  - spark=1.5  
  - pip  
  - pip:  
    - flask-migrate  
  - bar=1.4
```

```
$ conda env create  
$ source activate myenv
```

Manage multiple environments

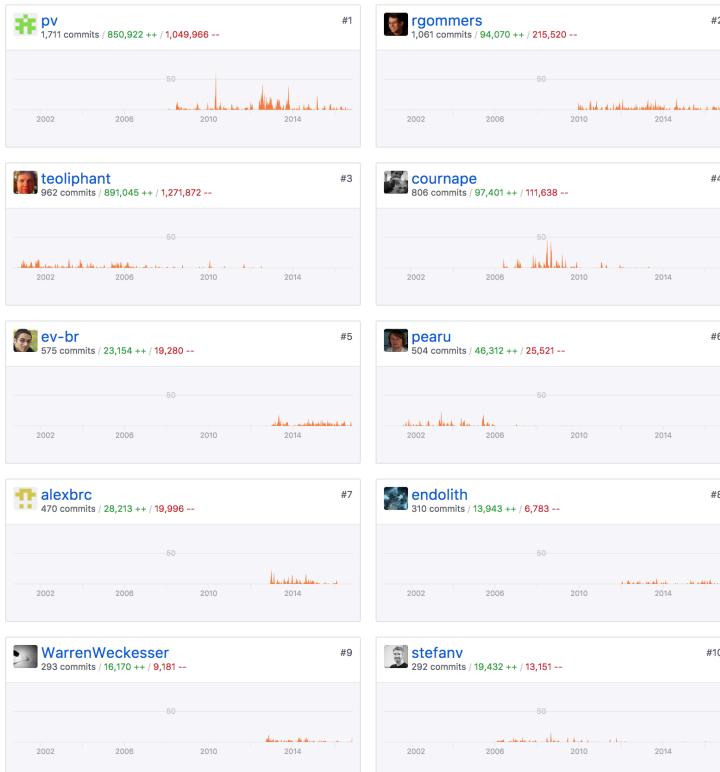


# Python Ecosystem

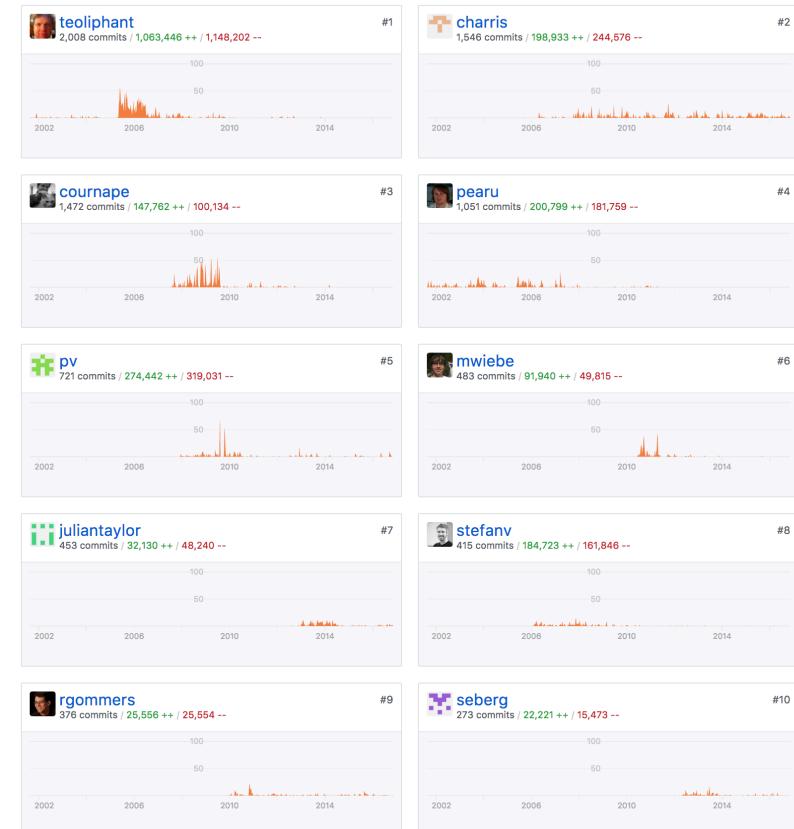


# An impressive community effort

SciPy ~ 478 contributors



NumPy ~ 509 contributors



# NumPy

```
from math import sin, pi

def sinc(x):
    if x == 0:
        return 1.0
    else:
        pix = pi*x
        return sin(pix)/pix

def step(x):
    if x > 0:
        return 1.0
    elif x < 0:
        return 0.0
    else:
        return 0.5
```

functions.py

## Without NumPy

```
>>> import functions as f
>>> xval = [x/3.0 for x in
range(-10,10)]
>>> yval1 = [f.sinc(x) for x
in xval]
>>> yval2 = [f.step(x) for x
in xval]
```

Python is a great language but needed a way to operate quickly and cleanly over multi-dimensional arrays.

```
from numpy import sin, pi
from numpy import vectorize
import functions as f

vsinc = vectorize(f.sinc)
def sinc(x):
    pix = pi*x
    val = sin(pix)/pix
    val[x==0] = 1.0
    return val

vstep = vectorize(f.step)
def step(x):
    y = x*0.0
    y[x>0] = 1
    y[x==0] = 0.5
    return y
```

## functions2.py

## With NumPy

```
>>> import functions2 as f
>>> from numpy import *
>>> x = r_[-10:10]/3.0
>>> y1 = f.sinc(x)
>>> y2 = f.step(x)
```

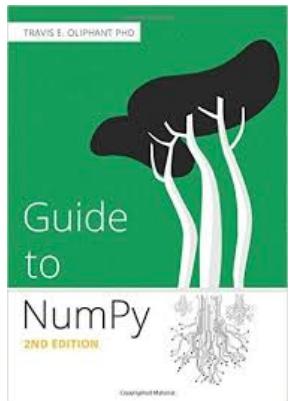
Offers N-D array, element-by-element functions, and basic random numbers, linear algebra, and FFT capability for Python

<http://numpy.org>

Fiscally sponsored by NumFOCUS

# NumPy: an Array Extension of Python

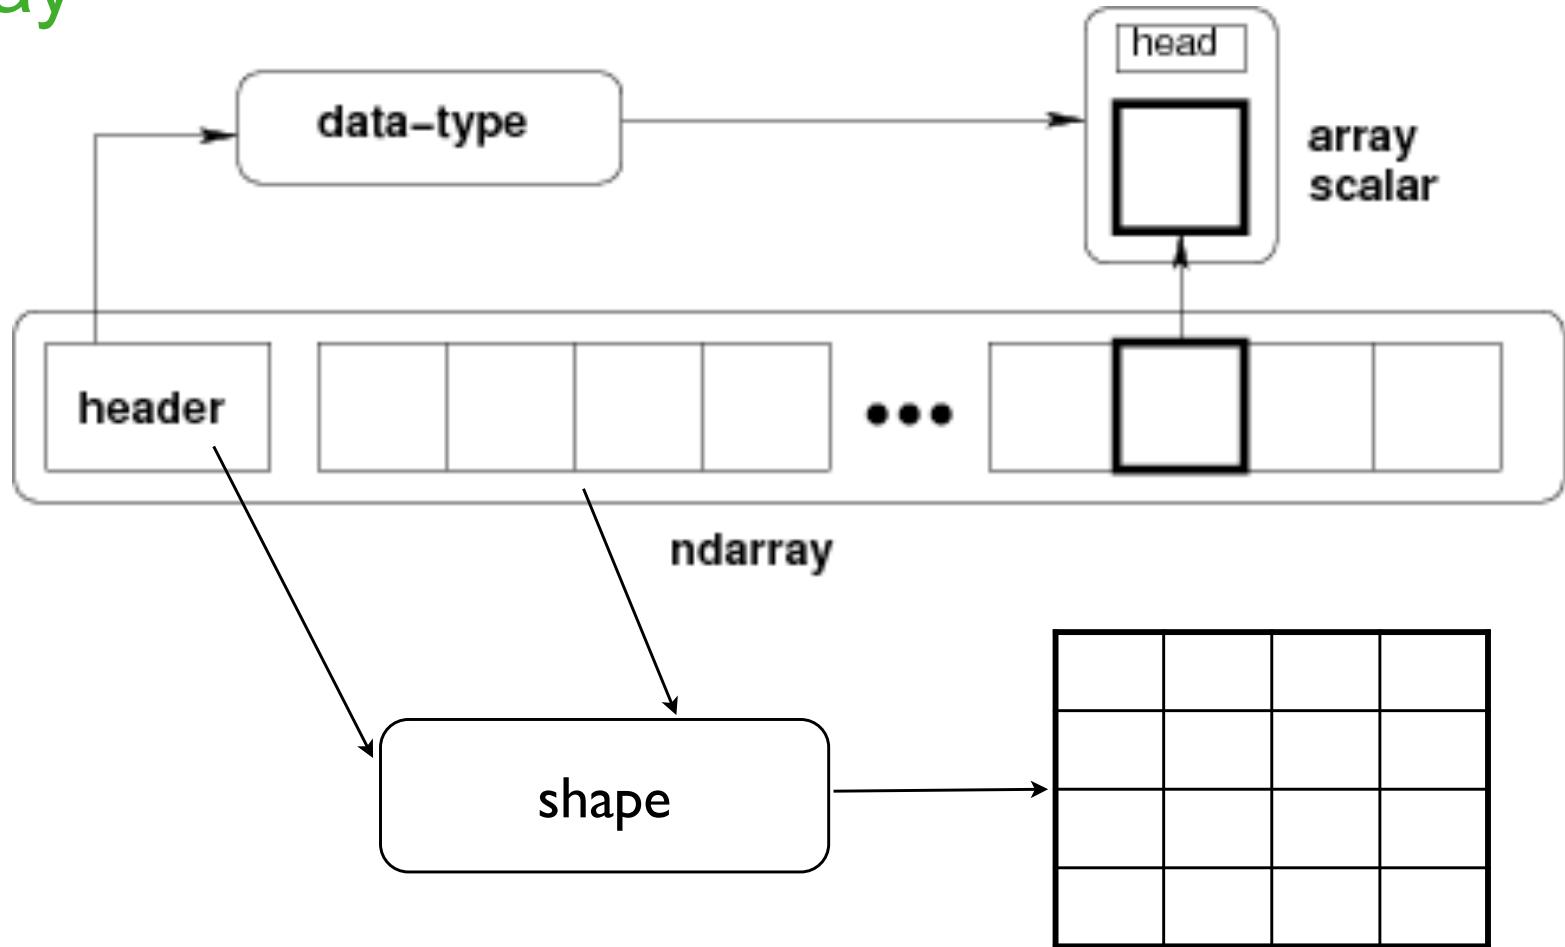
- Data: the array object
  - slicing and shaping
  - data-type map to Bytes
- Fast Math (ufuncs):
  - vectorization
  - broadcasting
  - aggregations



# NumPy Array

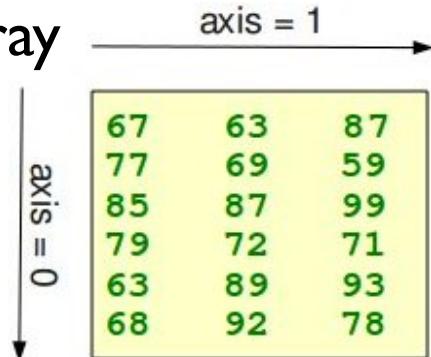
## Key Attributes

- dtype
- shape
- ndim
- strides
- data



# NumPy Examples

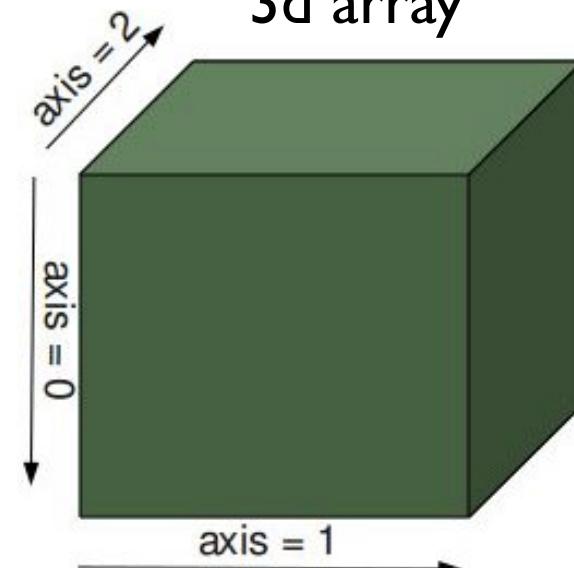
2d array



```
1 import numpy as np  
2  
3 x = np.array([[67, 63, 87],  
4 [77, 69, 59],  
5 [85, 87, 89],  
6 [79, 72, 71],  
7 [63, 89, 93],  
8 [68, 92, 78]])  
9 print x.sum(axis=0), x.sum(axis=1)
```

```
[439 472 477]  
[217 205 261 222 245 238]
```

3d array



```
12 y = 3*np.random.randn(10,20,30)+10  
13 print y.mean(), y.std()
```

9.98330639789 2.96677717122

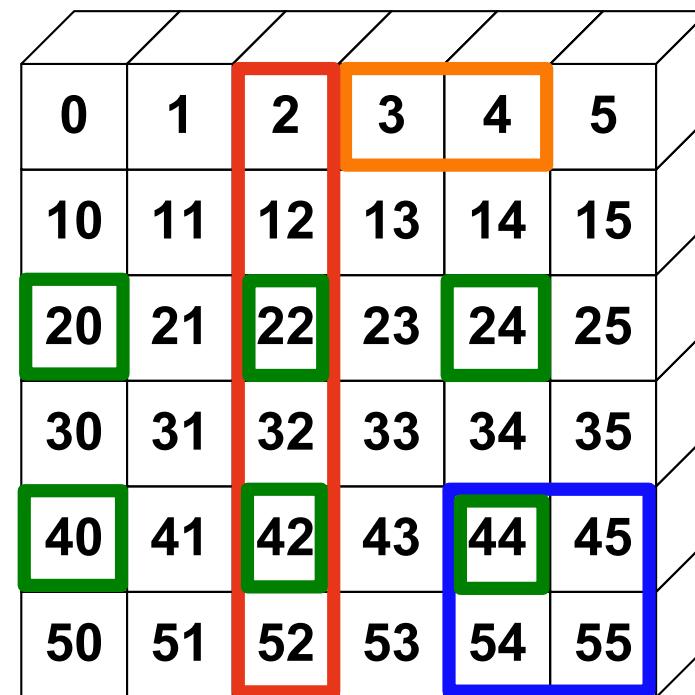
# NumPy Slicing (Selection)

```
>>> a[0,3:5]  
array([3, 4])
```

```
>>> a[4:,4:]  
array([[44, 45],  
      [54, 55]])
```

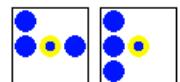
```
>>> a[:,2]  
array([2,12,22,32,42,52])
```

```
>>> a[2::2,:2]  
array([[20, 22, 24],  
      [40, 42, 44]])
```

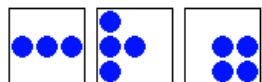


# Conway's game of Life

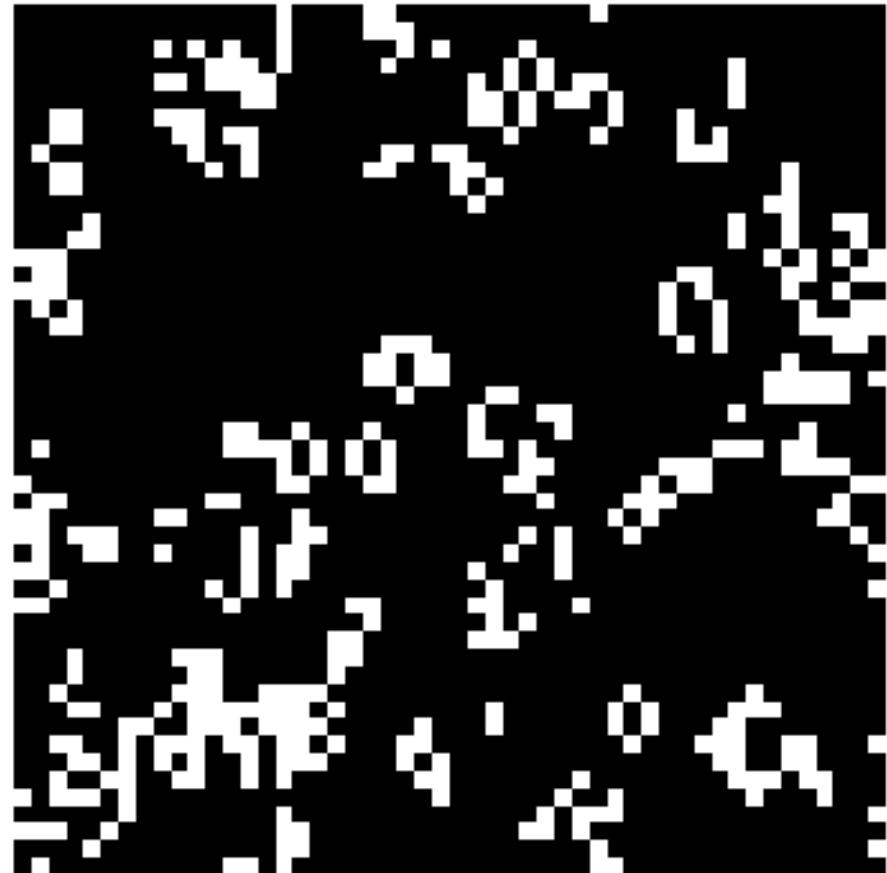
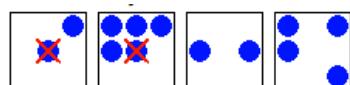
- Dead cell with exactly 3 live neighbors will come to life



- A live cell with 2 or 3 neighbors will survive



- With too few or too many neighbors, the cell dies



# Conway's Game of Life in NumPy

## Initialization

```
size = 100
GRID = (rand(size,size) > 0.5).astype(uint8)
# The world is round
indx = r_[0:size]
up = roll(indx, -1)
down = roll(indx, 1)
```

## Update Step

```
neighbors = GRID[up,:] + GRID[down,:] + GRID[:,up] + GRID[:,down] + \
GRID[ix_(up,up)] + GRID[ix_(up,down)] + GRID[ix_(down,up)] + GRID[ix_(down,down)]
GRID = (neighbors == 3) | (GRID & (neighbors==2))
```

# Summary

- Provides foundational N-dimensional array composed of homogeneous elements of a particular “dtype”
- The dtype of the elements is extensive (but not very extensible)
- Arrays can be sliced and diced with simple syntax to provide easy manipulation and selection.
- Provides fast and powerful math, statistics, and linear algebra functions that operate over arrays.
- Utilities for sorting, reading and writing data also provided.

# SciPy

# SciPy

“Distribution of Python Numerical Tools masquerading as a Library”

Name	Description
cluster	KMeans and Vector Quantization
fftpack	Discrete Fourier Transform
integrate	Numerical Integration
interpolate	Interpolation routines
io	Data Input and Output
linalg	Fast Linear algebra
misc	Utilities
ndimage	N-dimensional Image processing

Name	Description
odr	Orthogonal Distance Regression
optimize	Constrained and Unconstrained Optimization
signal	Signal Processing Tools
sparse	Sparse Matrices and Algebra
spatial	Spatial Data Structures and Algorithms
special	Special functions (e.g. Bessel)
stats	Statistical Functions and Distributions

## scipy.stats --- Continuous Distributions

94 continuous distributions!

Methods on all Random Variable Objects.

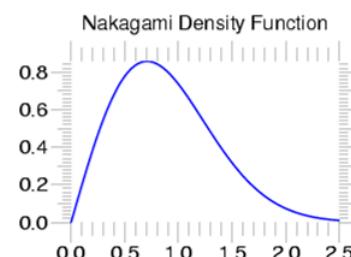
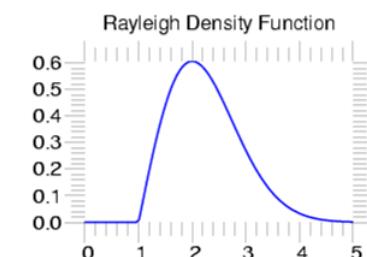
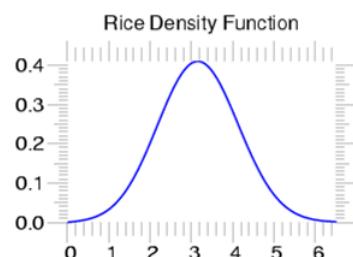
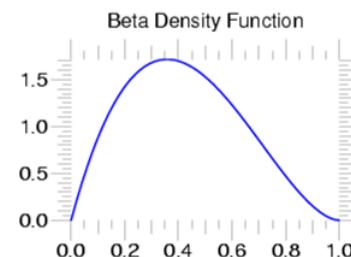
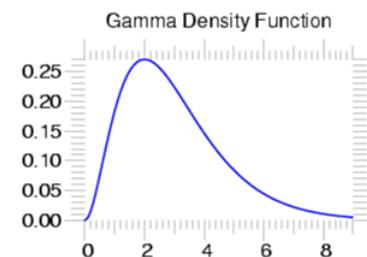
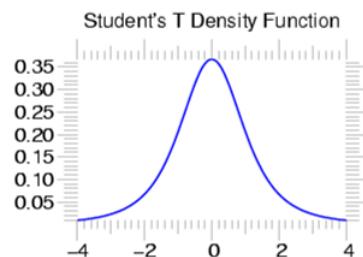
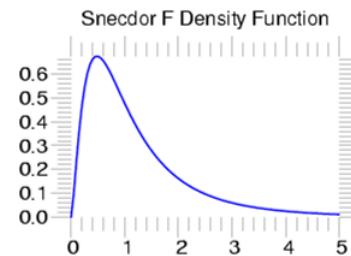
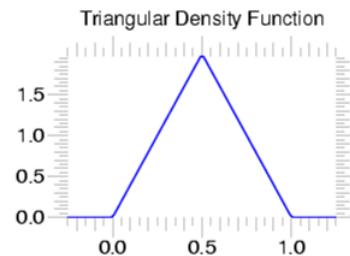
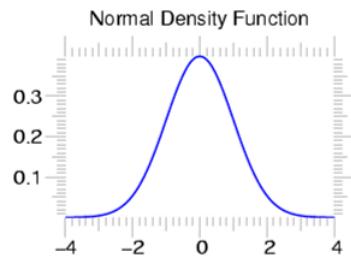
`pdf`

`cdf`

`rvs`

`ppf`

`stats`



## scipy.stats --- Discrete Distributions

13 standard discrete distributions (plus any arbitrary finite RV)

### Methods

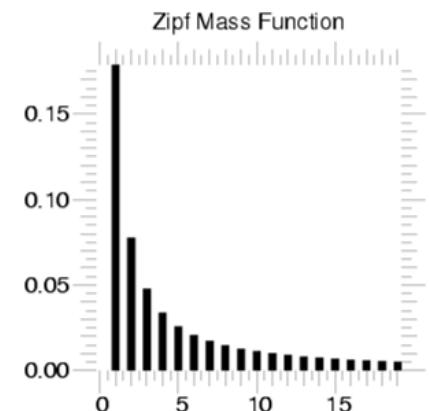
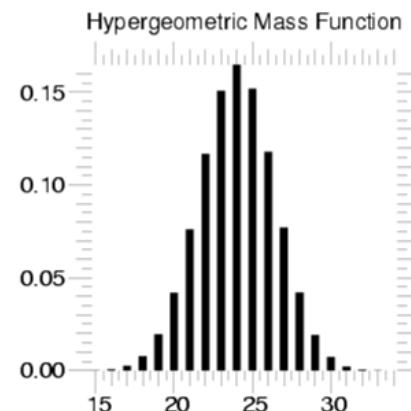
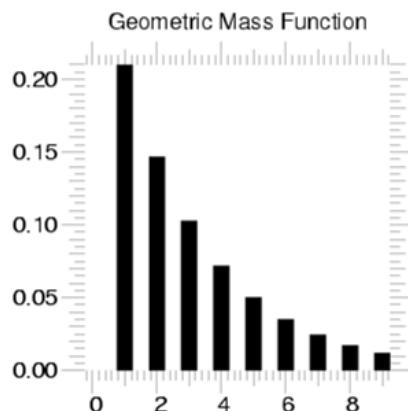
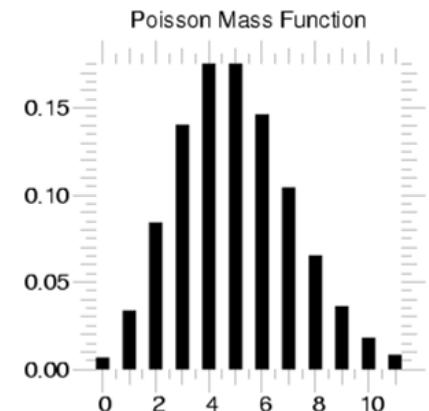
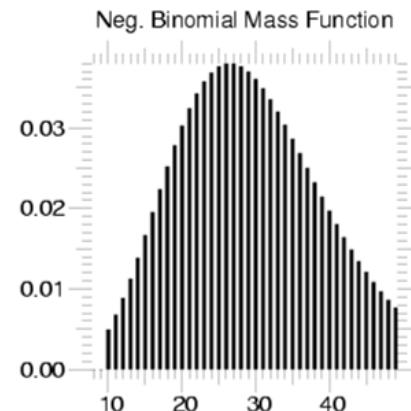
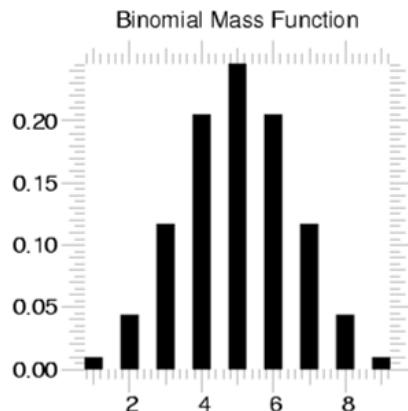
`pdf`

`cdf`

`rvs`

`ppf`

`stats`



# Matplotlib

# matplotlib

a powerful plotting engine

```
import numpy as np
import matplotlib.mlab as mlab
import matplotlib.pyplot as plt

np.random.seed(0)

# example data
mu = 100 # mean of distribution
sigma = 15 # standard deviation of distribution
x = mu + sigma * np.random.randn(437)

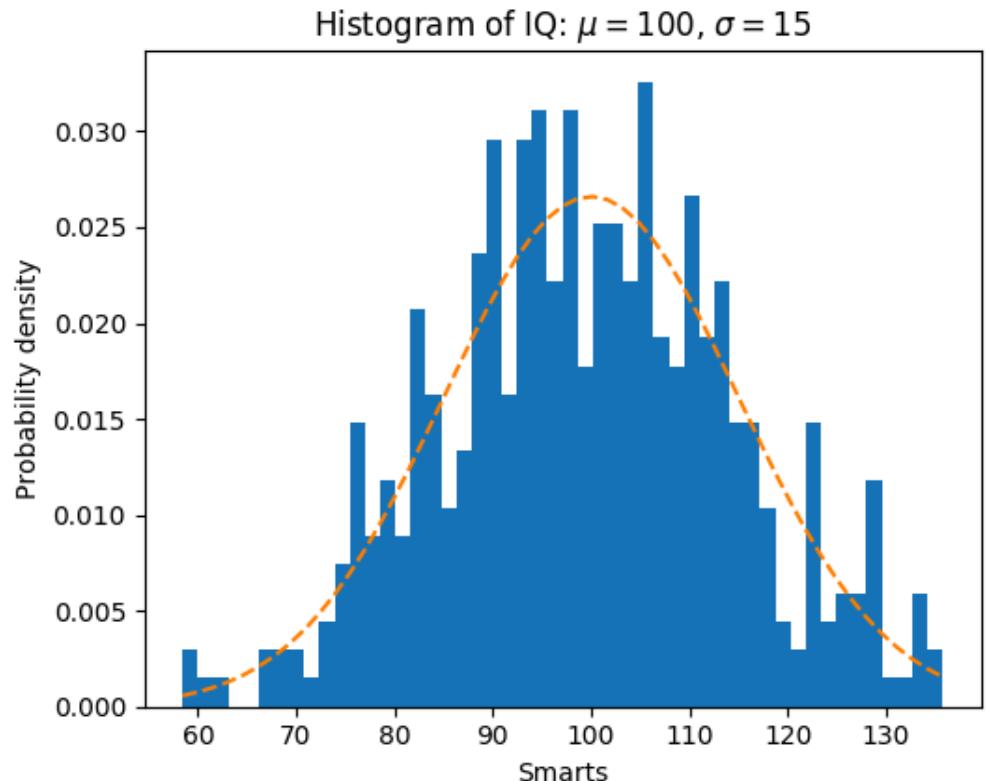
num_bins = 50

fig, ax = plt.subplots()

# the histogram of the data
n, bins, patches = ax.hist(x, num_bins, normed=1)

# add a 'best fit' line
y = mlab.normpdf(bins, mu, sigma)
ax.plot(bins, y, '--')
ax.set_xlabel('Smarts')
ax.set_ylabel('Probability density')
ax.set_title(r'Histogram of IQ: $\mu=100$', r'$\sigma=15$')

# Tweak spacing to prevent clipping of ylabel
fig.tight_layout()
plt.show()
```



# matplotlib

```
import matplotlib.pyplot as plt
import scipy.misc as misc

im = misc.face()
ax = plt.imshow(im)

plt.title('Racoon Face of size %d x %d' % im.shape[:2])

plt.savefig('face.png')
```

a powerful plotting engine



# matplotlib

```
import matplotlib.pyplot as plt
import numpy as np

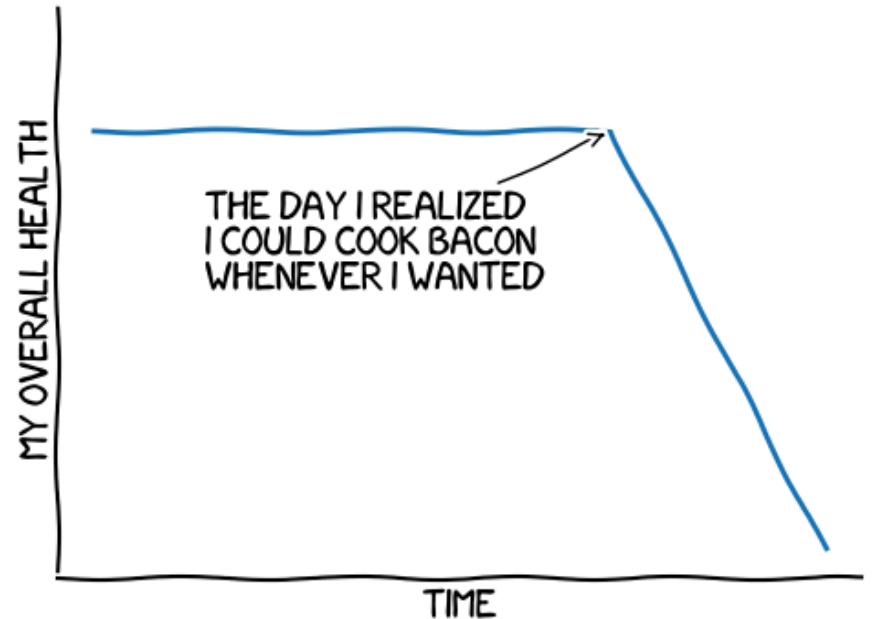
with plt.xkcd():
    fig = plt.figure()
    ax = fig.add_axes((0.1, 0.2, 0.8, 0.7))
    ax.spines['right'].set_color('none')
    ax.spines['top'].set_color('none')
    plt.xticks([])
    plt.yticks([])
    ax.set_ylim([-30, 10])
    data = np.ones(100)
    data[70:] -= np.arange(30)

    plt.annotate('THE DAY I REALIZED\nI COULD  
COOK BACON\nWHENEVER I WANTED',
                 xy=(70, 1),
                 arrowprops=dict(arrowstyle='->'),
                 xytext=(15, -10))

    plt.plot(data)

    plt.xlabel('time')
    plt.ylabel('my overall health')
    fig.text(0.5, 0.05, '"Stove Ownership"',
             from xkcd by Randall Monroe',
             ha='center')
```

a powerful plotting engine



"STOVE OWNERSHIP" FROM XKCD BY RANDALL MONROE

# matplotlib



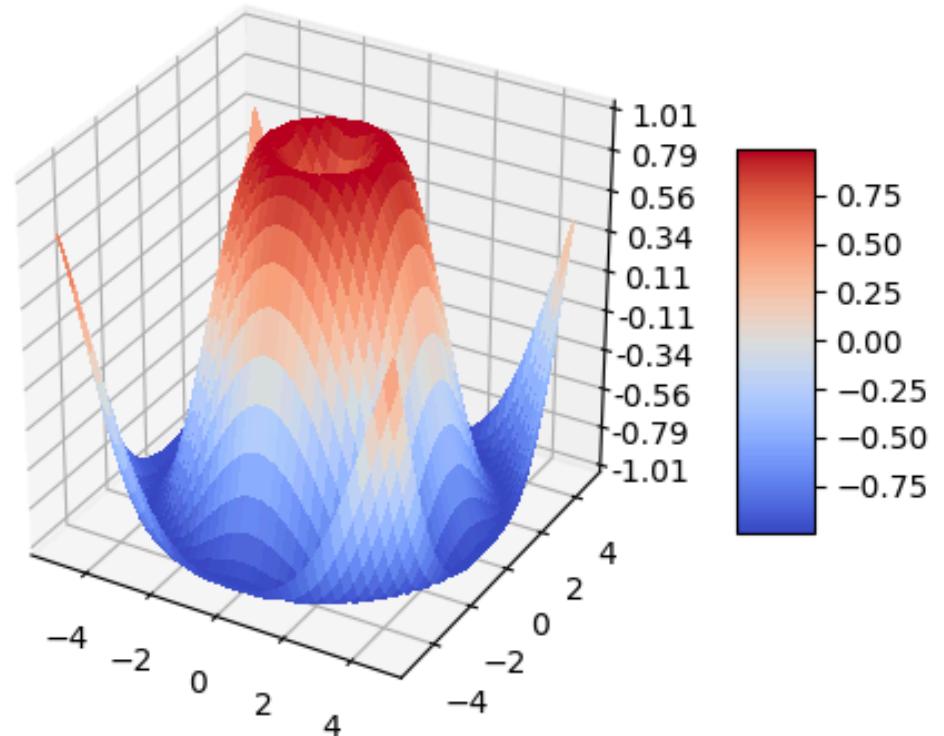
a powerful plotting engine

```
from mpl_toolkits.mplot3d import Axes3D
import matplotlib.pyplot as plt
from matplotlib import cm
from matplotlib.ticker import LinearLocator,
FormatStrFormatter
import numpy as np

fig = plt.figure()
ax = fig.gca(projection='3d')
# Make data.
X = np.arange(-5, 5, 0.25)
Y = np.arange(-5, 5, 0.25)
X, Y = np.meshgrid(X, Y)
R = np.sqrt(X**2 + Y**2)
Z = np.sin(R)
# Plot the surface.
surf = ax.plot_surface(X, Y, Z,
cmap=cm.coolwarm, linewidth=0,
antialiased=False)

# Customize the z axis.
ax.set_zlim(-1.01, 1.01)
ax.zaxis.set_major_locator(LinearLocator(10))
ax.zaxis.set_major_formatter(FormatStrFormatter(
'%.02f'))

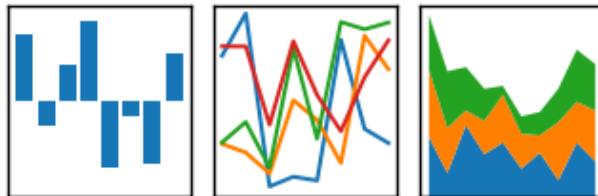
# Add a color bar which maps values to colors.
fig.colorbar(surf, shrink=0.5, aspect=5)
plt.show()
```



# Pandas

# pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$

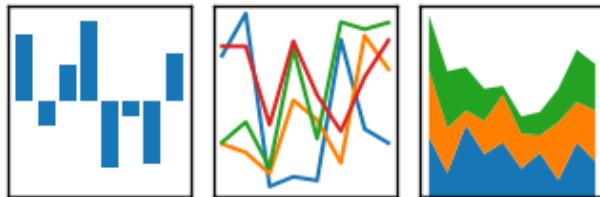


## Easy Data Wrangling

- Adds **indexes and labels** to 1-d and 2-d NumPy arrays (Series and DataFrame)
- Many convenience functions and methods to manipulate messy data-sets including **time-series**.
- Powerful indexing with automatic **data alignment**.
- Easy handling of **missing data**.
- Allows easy **joining and merging** Data Sets
- **Pivots** and reshaping (split-apply-combine)
- Powerful **group-by** operations with **summarization**
- Builtin visualization using labels and indexes

# pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$

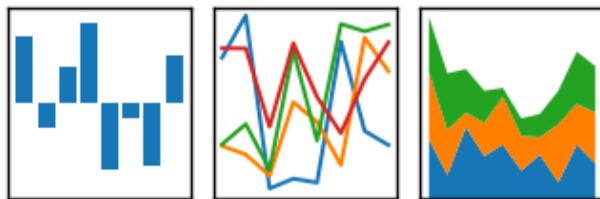


## Easy Data Wrangling

- Series Data Structure
  - built for **1-dimensional** series data
  - homogeneous data
  - Two arrays. One of data and another which is the index that can be a homogeneous array of any type like integers, objects, or date-times.
- DataFrame
  - built for **2-dimensional** collections of tabular data (think Excel sheet)
  - heterogeneous data comprised of multiple Series
  - includes an index column allowing sophisticated selection and alignment

# pandas

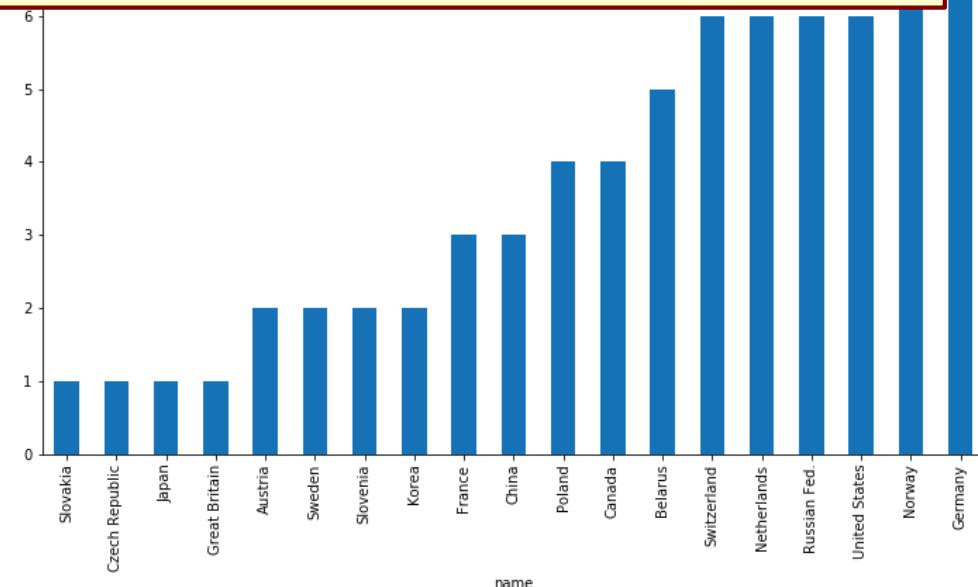
$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



## Easy Data Wrangling

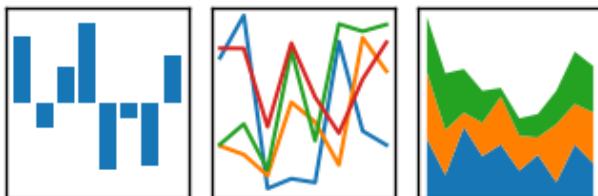
```
medals = pd.read_csv('data/medals.csv', index_col='name')
medals.head()
gold = medals['medal'] == 'gold'
won = medals['count'] > 0
medals.loc[gold & won, 'count'].sort_values().plot(kind='bar', figsize=(12,8))
```

	count	medal	country
name			
Australia	1	bronze	AUS
Australia	2	silver	AUS
Australia	0	gold	AUS
Austria	1	bronze	AUT
Austria	6	silver	AUT



# pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



## Easy Data Wrangling

```
google = pd.read_csv('data/goog.csv', index_col='Date', parse_dates=True)
google.info()
google.head()
google.describe()
```

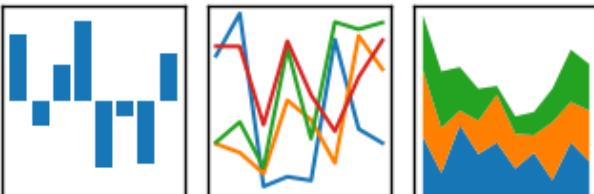
```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1257 entries, 2010-01-04 to 2014-12-31
Data columns (total 5 columns):
Open    1257 non-null float64
High    1257 non-null float64
Low     1257 non-null float64
Close   1257 non-null float64
Volume  194 non-null float64
dtypes: float64(5)
memory usage: 58.9 KB
```

	<b>Open</b>	<b>High</b>	<b>Low</b>	<b>Close</b>	<b>Volume</b>
<b>Date</b>					
<b>2010-01-04</b>	313.16	314.44	311.81	313.06	NaN
<b>2010-01-05</b>	313.28	313.61	310.46	311.68	NaN
<b>2010-01-06</b>	312.62	312.62	302.88	303.83	NaN
<b>2010-01-07</b>	304.40	304.70	296.03	296.75	NaN
<b>2010-01-08</b>	295.70	301.32	294.26	300.71	NaN

	<b>Open</b>	<b>High</b>	<b>Low</b>	<b>Close</b>	<b>Volume</b>
<b>count</b>	1257.000000	1257.000000	1257.000000	1257.000000	1.940000e+02
<b>mean</b>	375.275593	378.450247	372.132474	375.327064	1.937264e+06
<b>std</b>	115.684354	116.288827	114.935742	115.664301	9.842775e+05
<b>min</b>	218.940000	220.920000	216.600000	217.820000	7.040350e+05
<b>25%</b>	285.790000	288.760000	283.060000	285.450000	1.338451e+06
<b>50%</b>	318.330000	320.800000	315.180000	317.260000	1.684634e+06
<b>75%</b>	452.540000	456.020000	449.740000	452.830000	2.164369e+06
<b>max</b>	612.790000	613.830000	608.690000	609.470000	6.795393e+06

# pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



## Easy Data Wrangling

```
df = pd.read_excel("data/pbpython/salesfunnel.xlsx")
df.head()
table = pd.pivot_table(df,
                       index=["Manager", "Rep", "Product"],
                       values=["Price", "Quantity"],
                       aggfunc=[np.sum, np.mean])
```

	Account	Name	Rep	Manager	Product	Quantity	Price	Status
0	714466	Trantow-Barrows	Craig Booker	Debra Henley	CPU	1	30000	presented
1	714466	Trantow-Barrows	Craig Booker	Debra Henley	Software	1	10000	presented
2	714466	Trantow-Barrows	Craig Booker	Debra Henley	Maintenance	2	5000	pending
3	737550	Fritsch, Russel and Anderson	Craig Booker	Debra Henley	CPU	1	35000	declined
4	146832	Kiehn-Spinka	Daniel Hilton	Debra Henley	CPU	2	65000	won

Manager	Rep	Product	sum		mean	
			Price	Quantity	Price	Quantity
Debra Henley	Craig Booker	CPU	65000	2	32500.0	1.0
		Maintenance	5000	2	5000.0	2.0
		Software	10000	1	10000.0	1.0
	Daniel Hilton	CPU	105000	4	52500.0	2.0
		Software	10000	1	10000.0	1.0
	John Smith	CPU	35000	1	35000.0	1.0
		Maintenance	5000	2	5000.0	2.0
Fred Anderson	Cedric Moss	CPU	95000	3	47500.0	1.5
		Maintenance	5000	1	5000.0	1.0
		Software	10000	1	10000.0	1.0
	Wendy Yule	CPU	165000	7	82500.0	3.5
		Maintenance	7000	3	7000.0	3.0
		Monitor	5000	2	5000.0	2.0

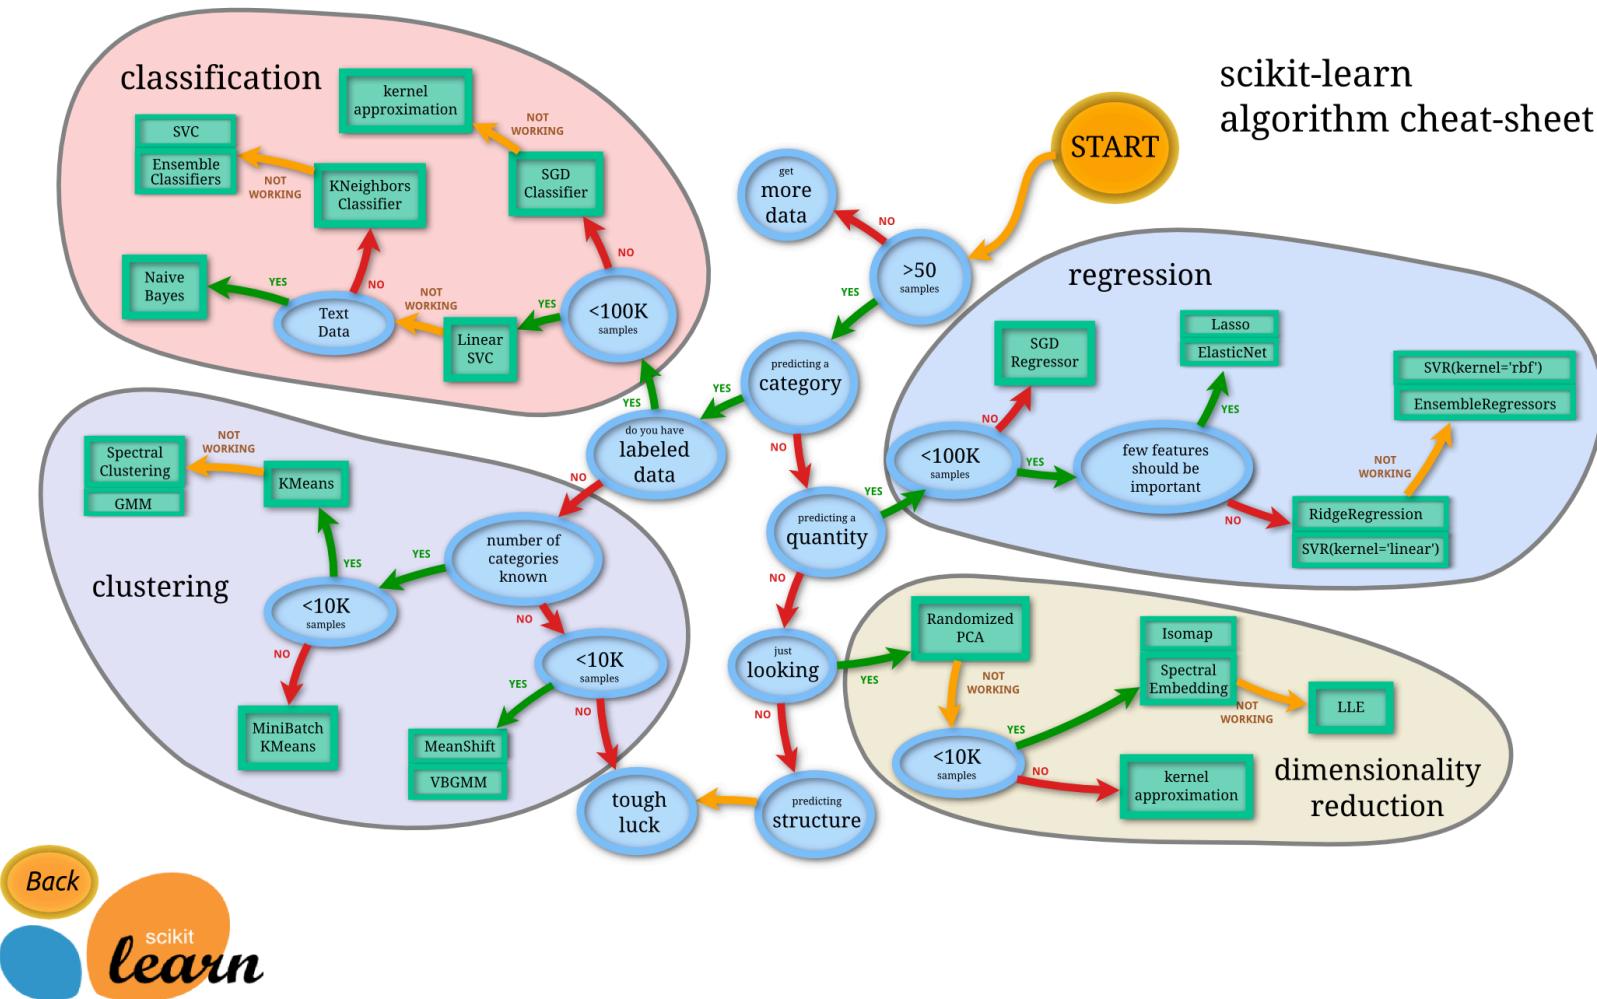
# Scikit-Learn



# Machine Learning made easy

- **Supervised Learning** — uses “labeled” data to train a model
  - Regression — predicted variable is continuous
  - Classification — predicted variable is discrete
- **Unsupervised Learning**
  - Clustering — discover categories in the data
  - Density Estimation — determine representation of data
  - Dimensionality Reduction — represent data with fewer variables or feature vectors
- Reinforcement Learning — “goal-oriented” learning (e.g. drive a car)
- **Deep Learning** — neural networks with many layers
- Semi-supervised Learning (use some labeled data for training)

## scikit-learn algorithm cheat-sheet





## Basic scikit-learn experience

### 1) Create or Load Data

```
>>> from sklearn import datasets  
>>> iris = datasets.load_iris()  
>>> digits = datasets.load_digits()
```

A Scikit-learn dataset is a “dictionary-like” object with input data stored as the `.data` attribute and labels stored as the `.target` attribute

`.data` attribute is always a 2D array (`n_samples, n_features`)

May need to extract features from raw data to produce data scikit-learn can use



## Basic scikit-learn experience

2) Choose Model (or Build Pipeline of Models)

```
>>> from sklearn import svm  
>>> clf = svm.SVC(gamma=0.001, C=100.)
```

Most models are model-families and have “hyper-parameters” that specify the specific model function. Good values for these can be found via grid-search and cross-validation (easy target for parallelization).

Here “gamma” and “C” are hyper-parameters.

Many choices of models: [http://scikit-learn.org/stable/supervised\\_learning.html](http://scikit-learn.org/stable/supervised_learning.html)



## Basic scikit-learn experience

### 3) Train the Model

```
>>> clf.fit(data[:-1], labels[:-1])
```

Models have a “fit” method which updates the parameters-to-be-estimated in the model in-place so that after fitting the model is “trained”

For validation and scoring you need to leave out some of the data to use later. cross-validation (e.g. k-fold) techniques can also be parallelized easily.

Here we “leave-one-out” (or n-fold)



## Basic scikit-learn experience

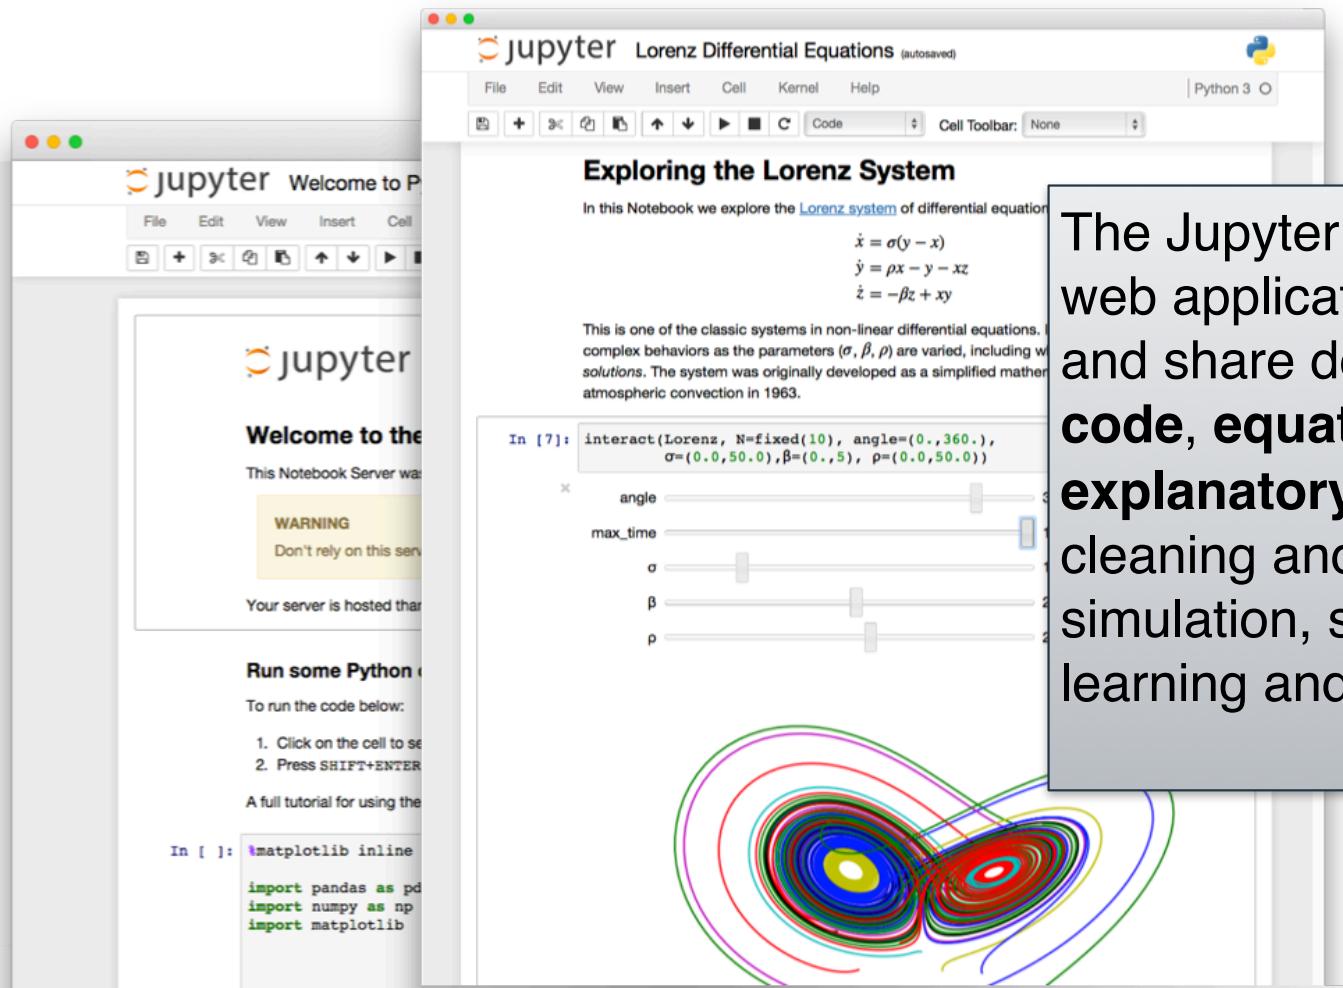
### 4) Predict new values

```
>>> clf.predict(data[-1:])
```

Prediction of new data uses the trained parameters of the model. Cross-validation can be used to understand how sensitive the model is to different partitions of the data.

```
>>> from sklearn.model_selection import cross_val_score
>>> scores = cross_val_score(clf, data, target, cv=10)
array([ 0.96...,  1. ...,  0.96...,  1.        ])
```

# Jupyter



The Jupyter Notebook is an open-source web application that allows you to create and share documents that contain live **code**, **equations**, **visualizations** and **explanatory text**. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, machine learning and much more.

# jupyter Flight Delays Modling (autosaved)



Logout

File Edit View Insert Cell Kernel Help | R O

[+][x][d][C] Markdown CellToolbar

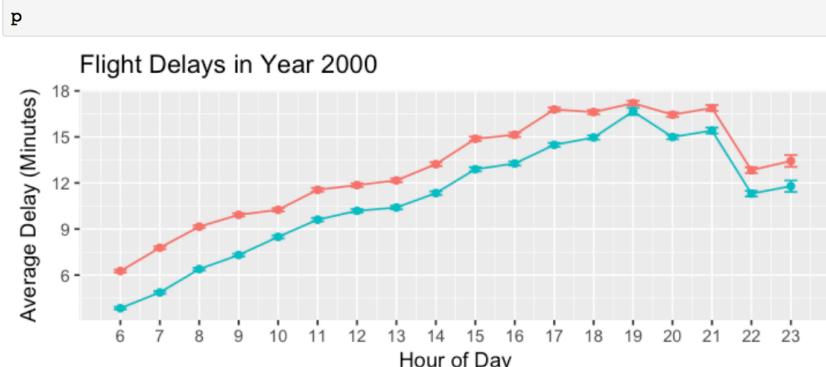
```
se=sqrt(var(newdelay,na.rm=TRUE)/length(na.omit(newdelay))
obs=length(na.omit(newdelay)))
```

In [7]: head(plot\_data, n=c(3L))

departure_hour	delay_type	mu	se	obs
6	Arrival Delay	6.258174	0.09168138	56981
6	Departure Delay	3.840884	0.08537210	57210
7	Arrival Delay	7.777764	0.09093139	62987

Using the R kernel allows us to make visualizations using ggplot natively.

In [10]: p=ggplot(plot\_data,aes(x=departure\_hour,y=mu,min=mu-se,max=mu+se,group=delay\_type))
 geom\_line() +
 geom\_point() +
 geom\_errorbar(width=.33) +
 scale\_x\_continuous(breaks=seq(6,23)) +
 labs(x="Hour of Day",y="Average Delay (Minutes)",title="Flight Delays in Year 2000",
 theme=legend.position="bottom") +
 scale\_color\_discrete(name="Delay Type")



# jupyter XKCD\_plots (autosaved)



Logout

[+][x][d][C] Markdown CellToolbar

In [5]:

```
np.random.seed(0)

ax = pylab.axes()

x = np.linspace(0, 10, 100)
ax.plot(x, np.sin(x) * np.exp(-0.1 * (x - 5) ** 2), 'b', lw=1, label='damped sine')
ax.plot(x, -np.cos(x) * np.exp(-0.1 * (x - 5) ** 2), 'r', lw=1, label='damped cosine')

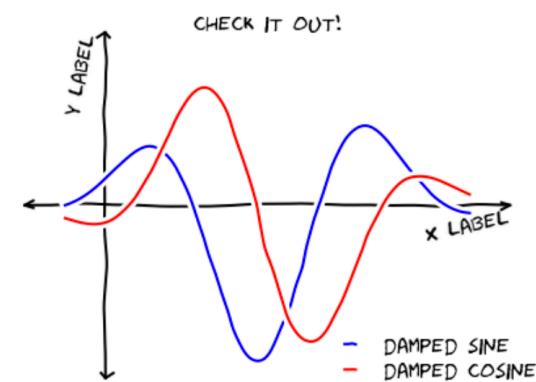
ax.set_title('check it out!')
ax.set_xlabel('x label')
ax.set_ylabel('y label')

ax.legend(loc='lower right')

ax.set_xlim(0, 10)
ax.set_ylim(-1.0, 1.0)

#XKCDify the axes -- this operates in-place
XKCDify(ax, xaxis_loc=0.0, yaxis_loc=1.0,
         xaxis_arrow='+', yaxis_arrow='+-',
         expand_axes=True)
```

Out[5]: <matplotlib.axes.AxesSubplot at 0x2fecbd0>



# Collaborative Executable Notebooks

ANACONDA Airline Industry Project Notebook Last Checkpoint: 03/01/2016 (unsaved changes)

File Edit View Insert Cell Kernel Help

Code CellToolbar i p

In [1]: 

```
from bokeh.io import output_notebook, push_notebook
from bokeh.models import ColumnDataSource, HoverTool
from bokeh.palettes import RdYlGn8
from bokeh.plotting import figure, show
from bokeh.resources import INLINE
import numpy as np
import pandas as pd

output_notebook(resources=INLINE)
```

BokehJS successfully loaded.

Data lineage

Flight delays

Flight delay data are loaded from csv and aggregated by airport and monthly totals. For visualization purposes green-to-red color map to describe delay frequency and airports are binned into size groups based on the number of flights.

In [2]: 

```
flights = pd.read_csv('data/airline_delay_causes.csv',
                       parse_dates={'period': [0,1]},
                       usecols=['year', 'month', 'airport', 'arr_flights', 'arr_delay'])

flights = flights.groupby(['airport', 'period']).sum()
flights['late_pct'] = flights.arr_delay / flights.arr_flights
flights['color'] = pd.qcut(flights.late_pct, q=8, labels=RdYlGn8)
flights['size'] = pd.cut(flights.arr_flights, bins=[0, 500, 5000, 50000], labels=[1, 2, 3, 4])
flights.head(3)
```

Out[2]:

airport	period	arr_flights	arr_delay	late_pct	color	size
2003-06-01		334	69	0.206587	#fee08b	5

Advanced notebook extensions

ANACONDA Airline Industry Project Notebook Last Checkpoint: 03/01/2016 (unsaved changes)

File Edit View Insert Cell Kernel Help

Code CellToolbar i p

US Airport Delays

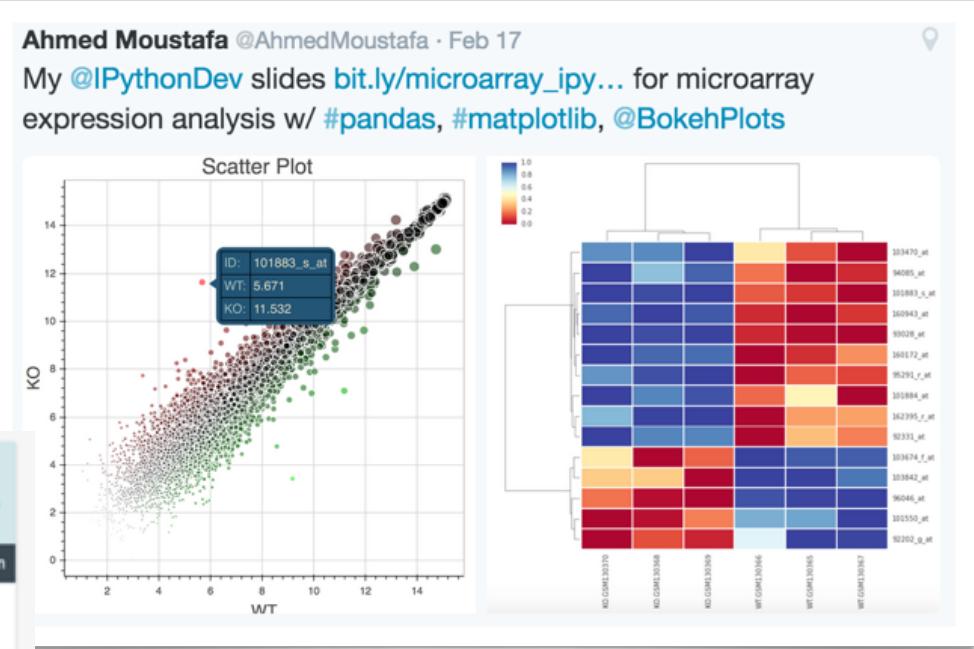
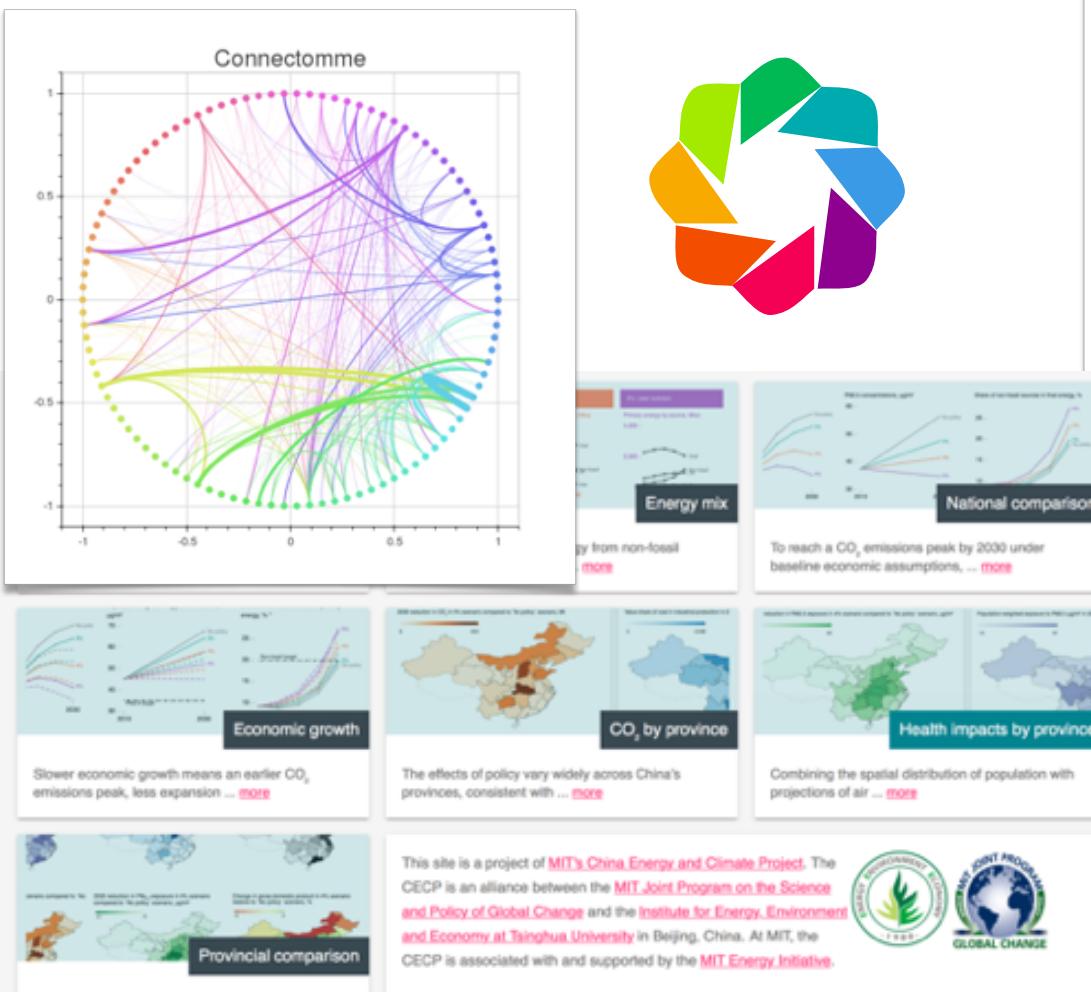
Latitude (deg)

Longitude (deg)

Interactive Visualizations

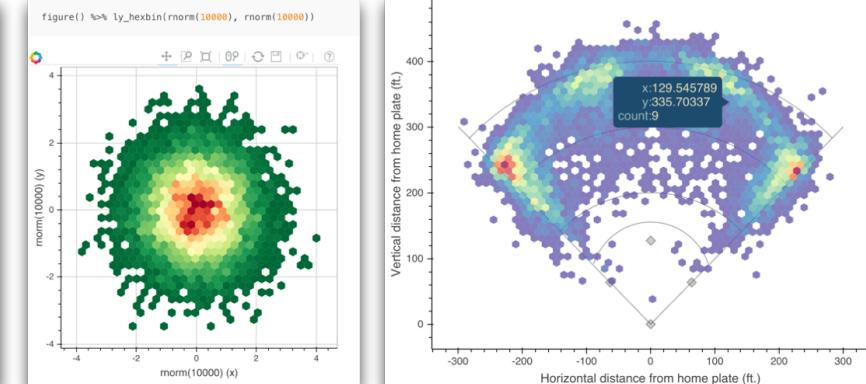
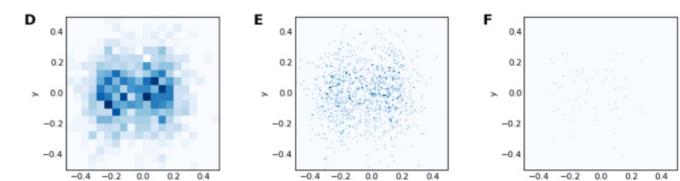
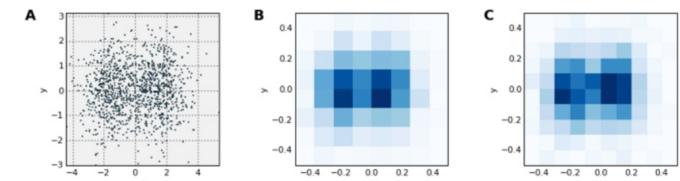
# Interactive Data Visualization Apps with Bokeh

# Interactive Data Visualization



- Interactive viz, widgets, and tools
- Versatile high level graphics
- Streaming, dynamic, large data
- Optimized for the browser
- No Javascript
- With or without a server

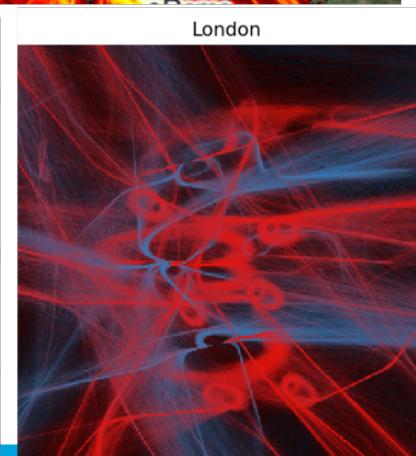
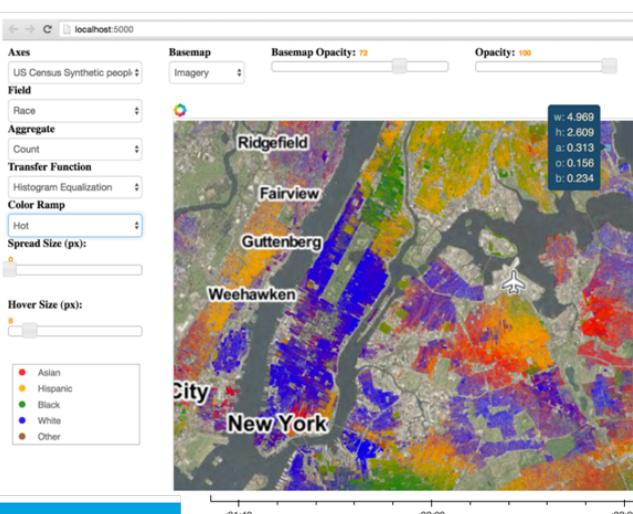
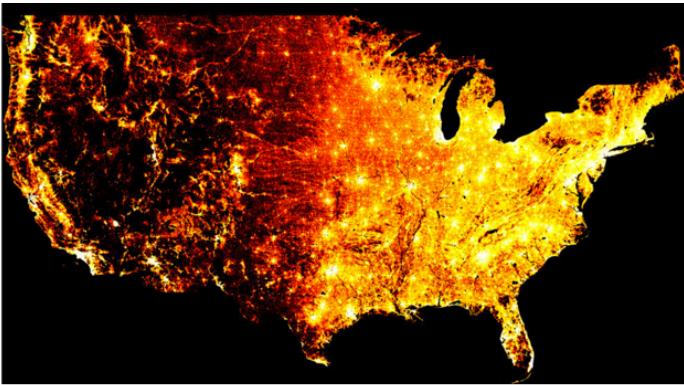
# Rapid Prototyping Visual Apps



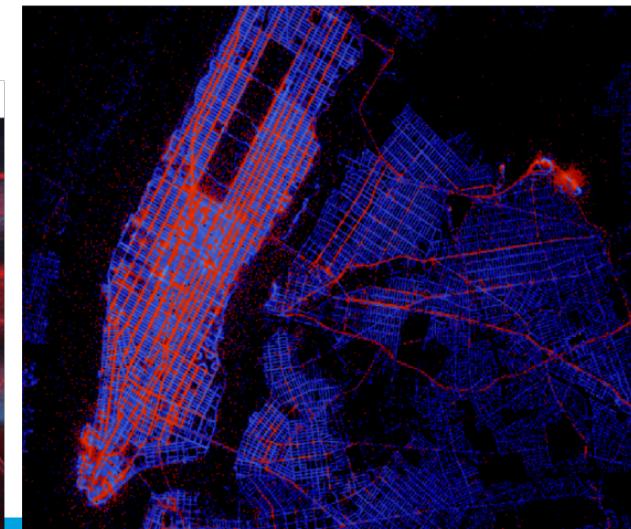
- Python interface
- R interface
- Smart plotting

# Plotting Billions of Points and Map Integration with Datashader

# Datashader: Rendering a Billion Points of Data

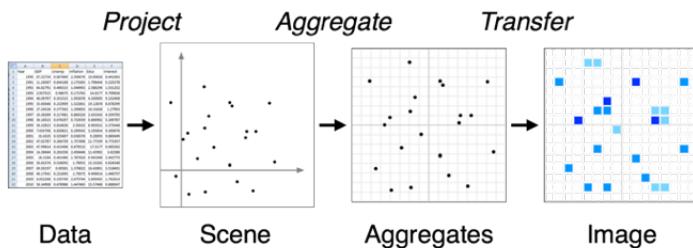


- datashader provides a fast, configurable visualization pipeline for faithfully revealing even very large datasets
- Each of these visualizations requires just a few lines of code and no magic numbers to adjust by trial and error.



# Datashader

## Configurable visualization pipeline:



## Simple code, with no magic numbers:

```
cvs = ds.Canvas(plot_width=900, plot_height=600)
agg = cvs.points(df, 'longitude', 'latitude', ds.count())
img = tf.shade(agg, cmap=hot, how='eq_hist')
```

## Powerful computations on the data:

E.g. show only top 10% highest-count pixels, or pixels where blacks outnumber whites:

```
agg2 = agg.where(agg>=np.percentile(agg,90))
agg3 = agg.where(agg.sel(race='w') < agg.sel(race='b'))
```

## High performance

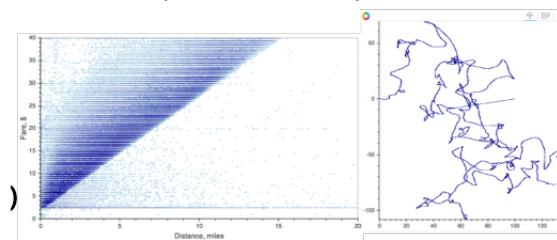
Only the first stage depends on the dataset size — visualizes arbitrarily large datasets, without downsampling, in your browser.

Fast on ordinary machines, and scales up easily for distributed, out of core processing

E.g. 10 million points in 0.5 seconds on a laptop, or 3 billion points in 10 seconds on a cluster.

## Flexible datatypes

- Scatterplots/heatmaps
- Time series
- Connected points (trajectories)
- Rasters (via Rasterio)



## Highly interactive:

Output can be embedded into Bokeh plots with pan,zoom,overlays,axes



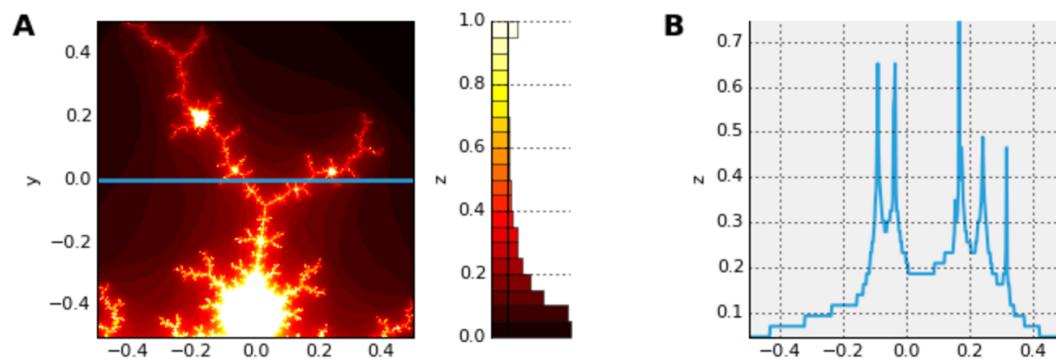
# Data Visualization and Applications made easy with Holoviews

# HoloViews: Stop plotting your data

HoloViews makes it simple to create beautiful interactive Bokeh or Matplotlib visualizations of complex data.

```
import numpy as np
import holoviews as hv
hv.notebook_extension('matplotlib')
fractal = hv.Image(np.load('mandelbrot.npy'))

((fractal * hv.HLine(y=0)).hist() + fractal.sample(y=0))
```



- Exploring data can be tedious if you use a plotting library directly, because you will need to specify details about your data (units, dimensions, names, etc.) every time you construct a new type of plot.
- With HoloViews, you instead annotate your data once, and then flexible plotting comes for free — HoloViews objects just display themselves, alone or in any combination.
- It's now easy to lay out subfigures, overlay traces, facet or animate a multidimensional dataset, sample or aggregate to reduce dimensionality, preserving the metadata each time so the results visualize themselves.

[https://anaconda.org/jbednar/nyc\\_taxi-paramnb/notebook](https://anaconda.org/jbednar/nyc_taxi-paramnb/notebook)

NYC Taxi trips, with [Datashader](#), [HoloViews](#), [GeoViews](#) and [ParamNB](#)

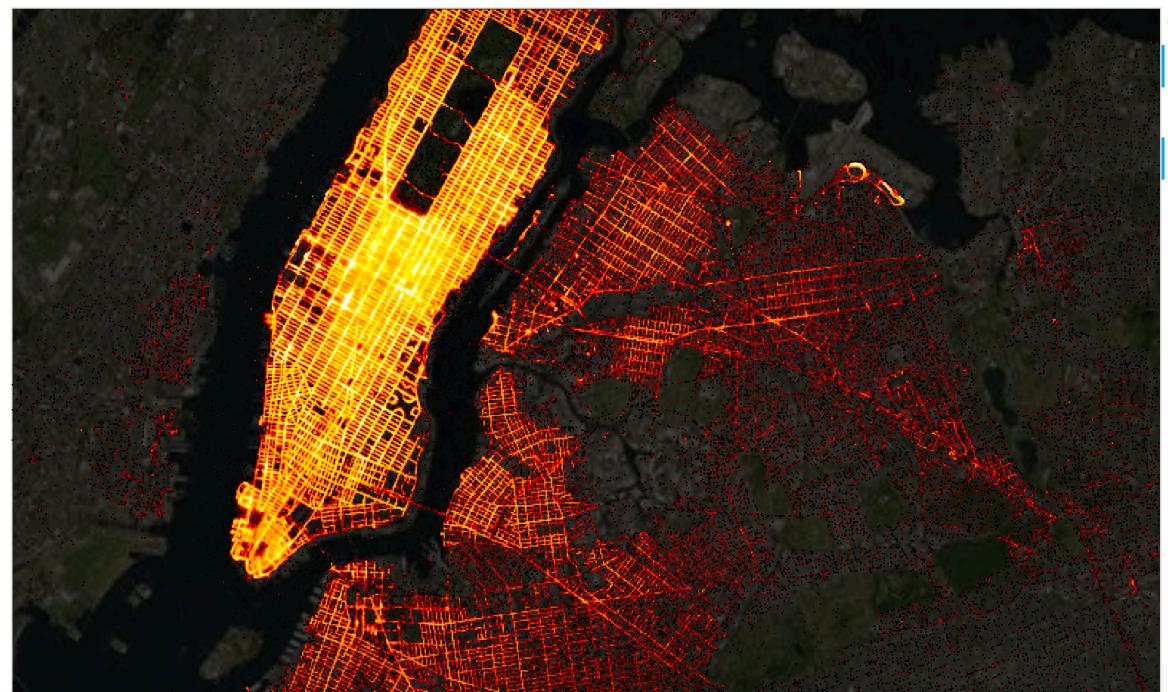
# Data Widgets and Applications from Jupyter Notebooks!

alpha  0.30

colormap

passengers

plot



```
tiles = gv.WMTSTileSource(url='https://server.arcgisonline.com/Ar
    'World_Imagery/MapServer/tile/{Z}/{Y}
tile_options = dict(width=800,height=475,xaxis=None,yaxis=None,bgcolor=
passenger_counts = sorted(df.passenger_count.unique().tolist())

class Options(hv.streams.Stream):
    alpha      = param.Magnitude(default=0.75, doc="Alpha value for th
    colormap   = param.ObjectSelector(default=cm["fire"], objects=cm.v
    plot       = param.ObjectSelector(default="pickup",   objects=[ "pi
    passengers = param.ObjectSelector(default=1,           objects=pass

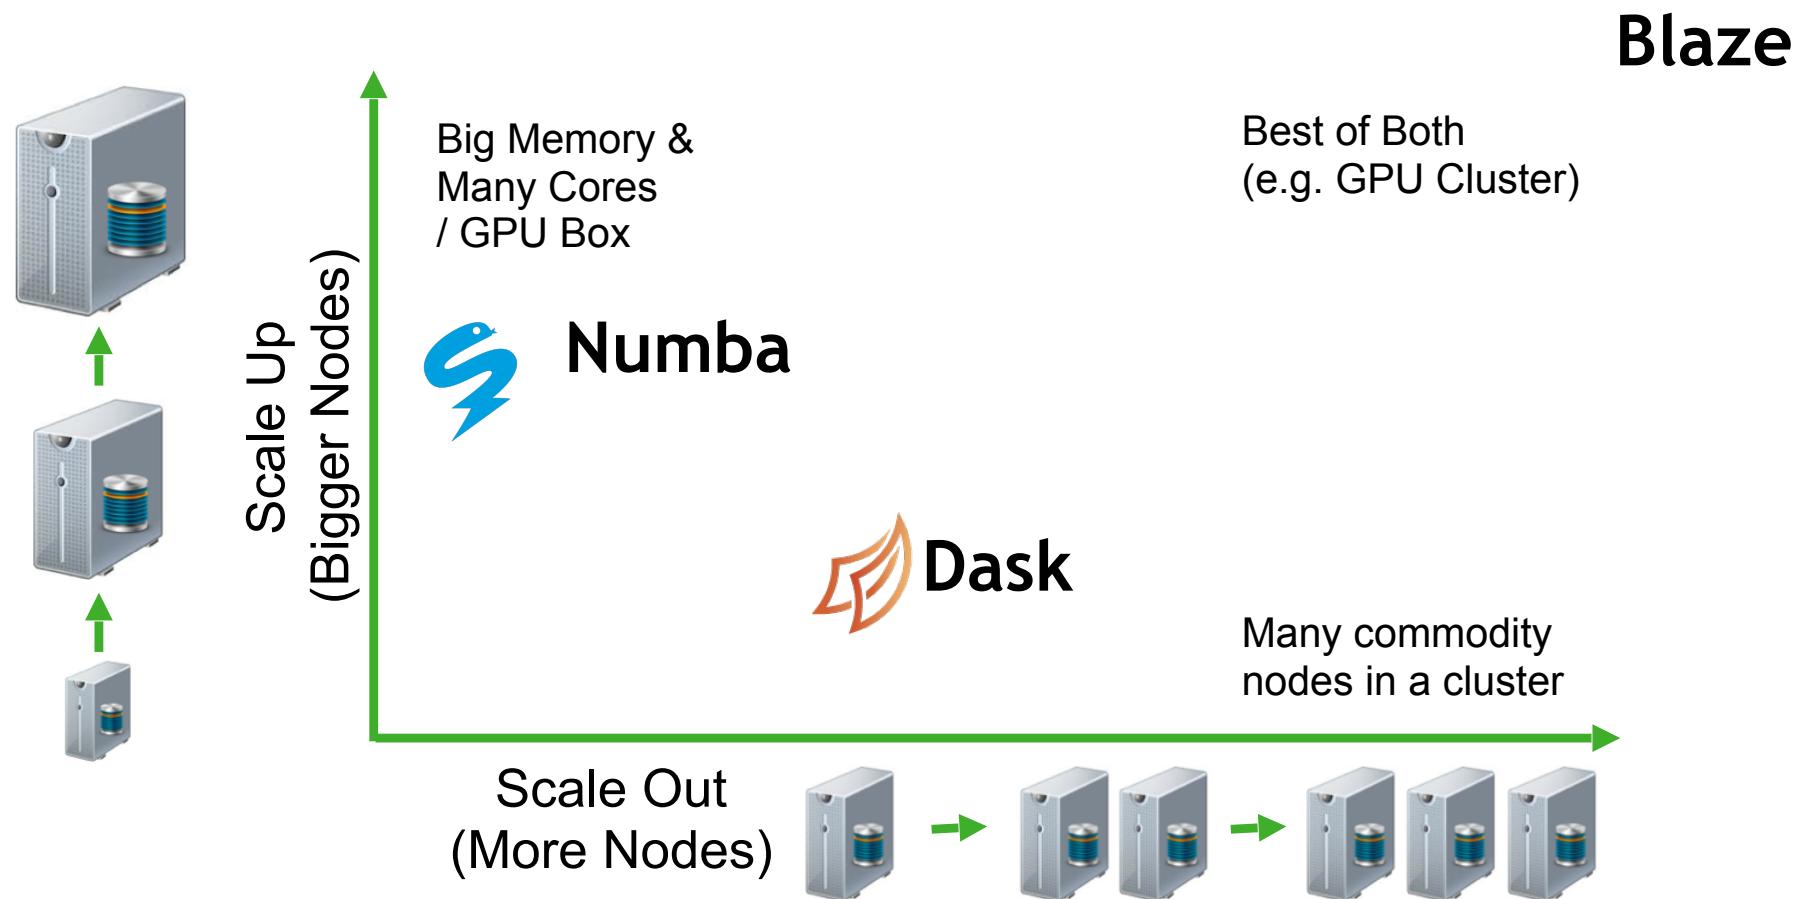
    def make_plot(self, x_range=None, y_range=None, **kwargs):
        map_tiles = tiles(style=dict(alpha=self.alpha), plot=tile_optio
            df_filt = df[df.passenger_count==self.passengers]
            points = hv.Points(gv.Dataset(df_filt, kdims=[self.plot+'_x', s
            taxi_trips = datashade(points, width=800, height=475, x_samplin
                cmap=self.colormap, element_type=gv.Image,
                dynamic=False, x_range=x_range, y_range=y_range)

        return map_tiles * taxi_trips

selector = Options(name="")
paramnb.Widgets(selector, callback=selector.update)
hv.DynamicMap(selector.make_plot, kdims=[], streams=[selector, RangeXY]
```

# Scaling Up and Out with Numba and Dask

# Scale Up vs Scale Out



# Scaling Up! Optimized Python with JIT compilation from Numba

# Numba



- Started in 2012
- Release 31 (0.31) in February
- Version 1.0 coming in 2017
- Particularly suited to Numeric computing
- Lots of features!
- Ahead of Time Compilation
- Wide community adoption and use

```
conda install numba
```

# Example: Filter an array

In [87]:

```
@jit(nopython=True)
def nan_compact(x):
    out = np.empty_like(x)
    out_index = 0
    for element in x:
        if not np.isnan(element):
            out[out_index] = element
            out_index += 1
    return out[:out_index]
```

Annotations:

- `@jit(nopython=True)` → Numba decorator  
(`nopython=True` not required)
- `out = np.empty_like(x)` → Array Allocation
- `for element in x:` → Looping over ndarray `x` as an iterator
- `if not np.isnan(element):` → Using numpy math functions
- `return out[:out_index]` → Returning a slice of the array

In [88]:

```
a = np.random.uniform(size=10000)
a[a < 0.2] = np.nan
np.testing.assert_equal(nan_compact(a), a[~np.isnan(a)])
```

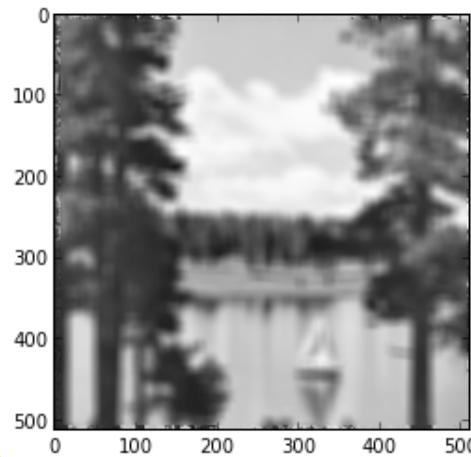
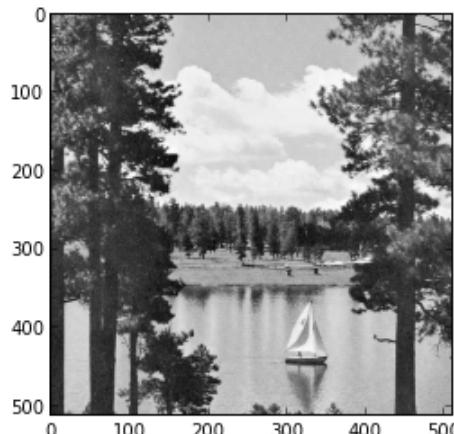
In [89]:

```
%timeit a[~np.isnan(a)]
%timeit nan_compact(a)

10000 loops, best of 3: 52 µs per loop
100000 loops, best of 3: 19.6 µs per loop
```

2.7x Speedup  
over NumPy!

# Image Processing



~1500x speed-up

```
@jit('void(f8[:, :], f8[:, :], f8[:, :])')
def filter(image, filt, output):
    M, N = image.shape
    m, n = filt.shape
    for i in range(m//2, M-m//2):
        for j in range(n//2, N-n//2):
            result = 0.0
            for k in range(m):
                for l in range(n):
                    result += image[i+k-m//2, j+l-n//2]*filt[k, l]
            output[i, j] = result
```

# Numba Compatibility

Does *not* replace the standard Python interpreter

OS	HW	SW
Windows (7 and later)	32 and 64-bit x86 CPUs	Python 2 and 3
OS X (10.9 and later)	CUDA-capable NVIDIA GPUs	NumPy 1.7 through 1.11
Linux (~RHEL 5 and later)	HSA-capable AMD GPUs	

# Scaling Out with Dask (integrates with but doesn't depend on Hadoop)

# Dask

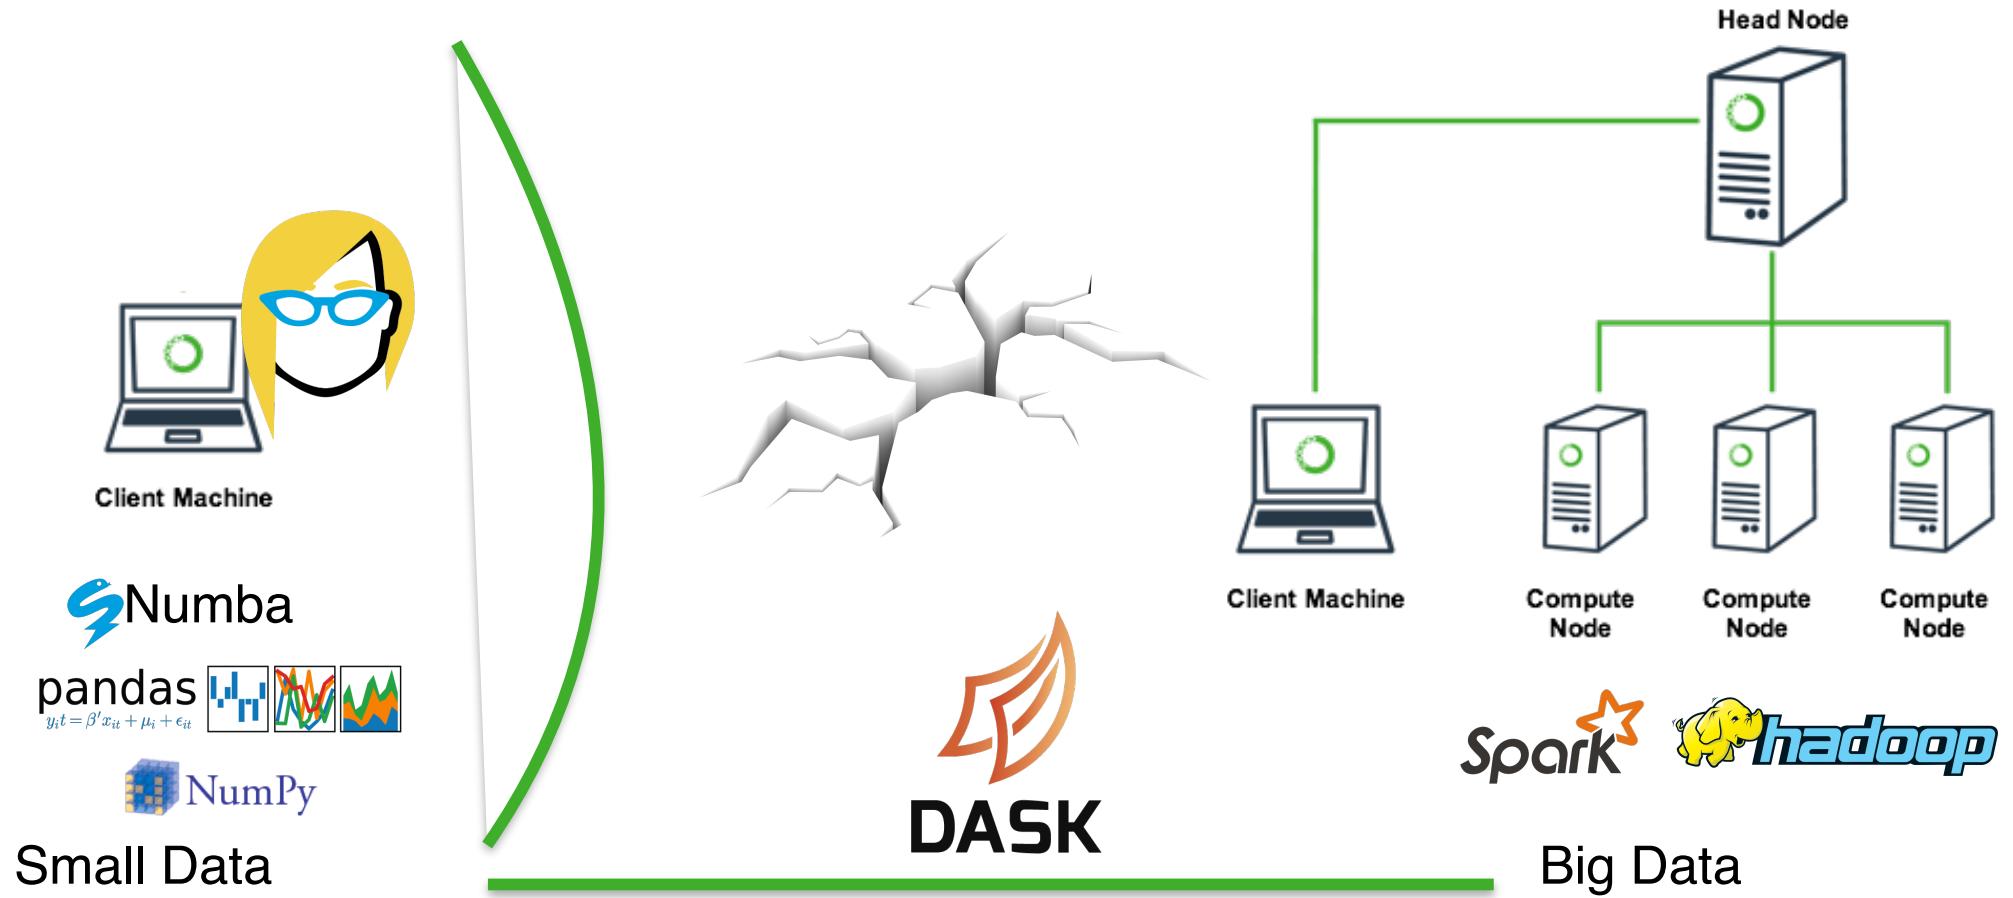


- Started as part of Blaze in early 2014.
- General parallel programming engine
- Flexible and therefore highly suited for
  - Commodity Clusters
  - Advanced Algorithms
- Wide community adoption and use

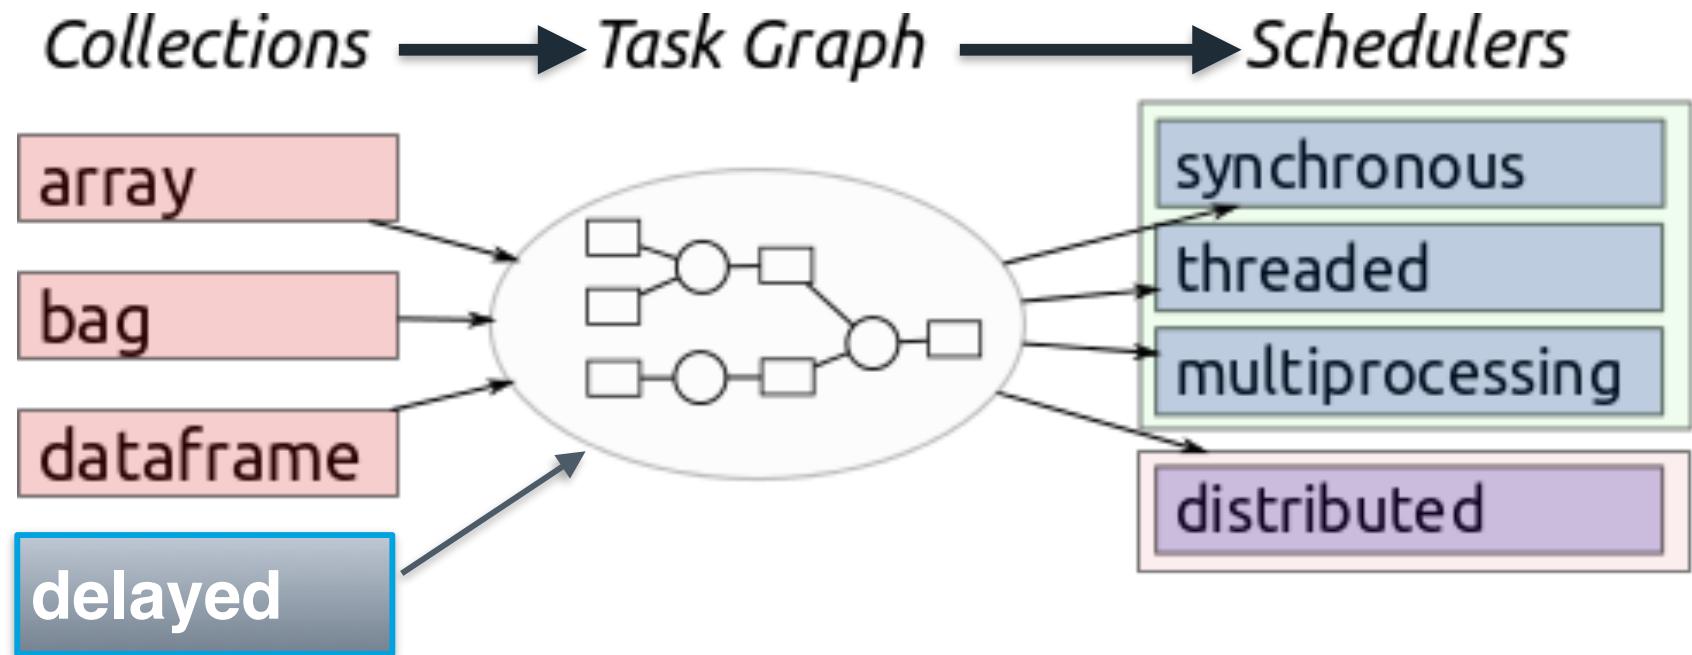
```
conda install -c conda-forge dask
```

```
pip install dask[complete] distributed --upgrade
```

# Moving from small data to big data

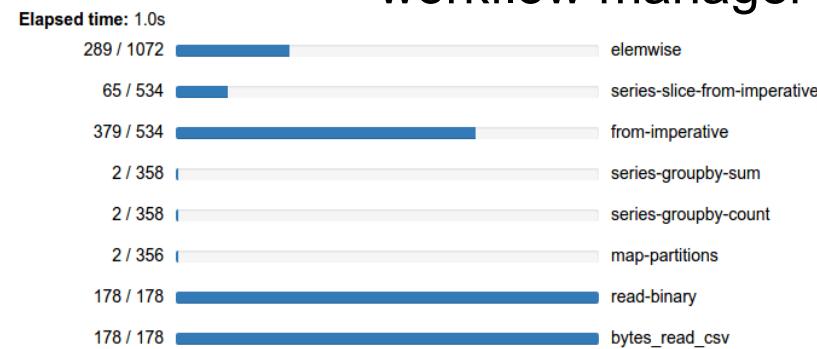
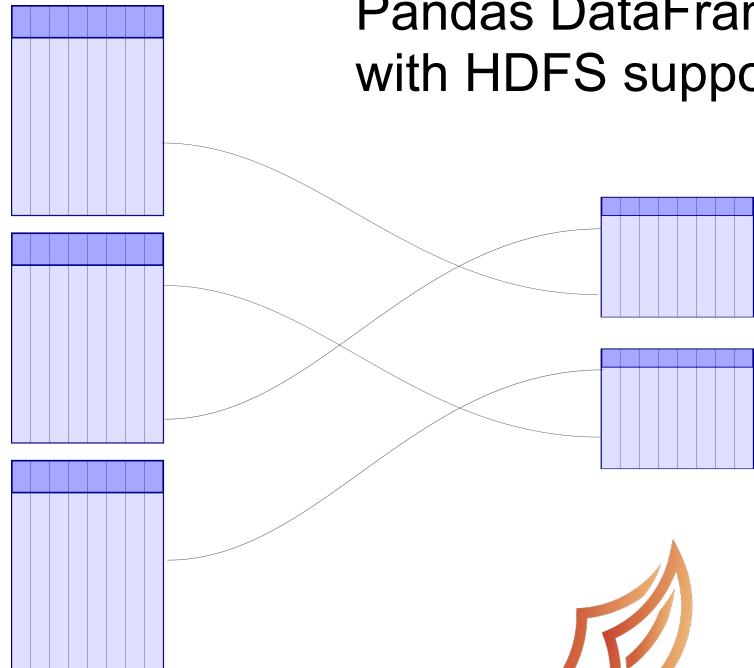


# Dask: From User Interaction to Execution

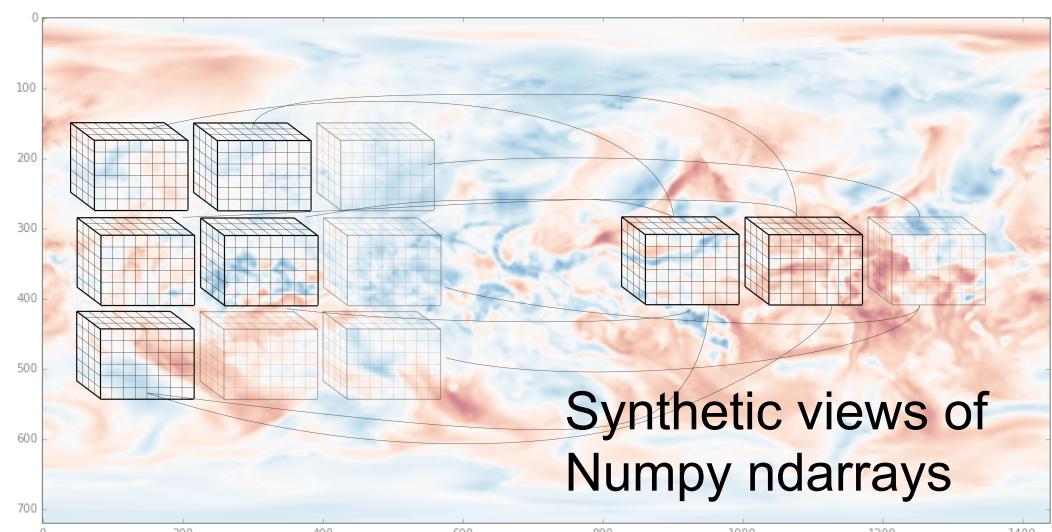
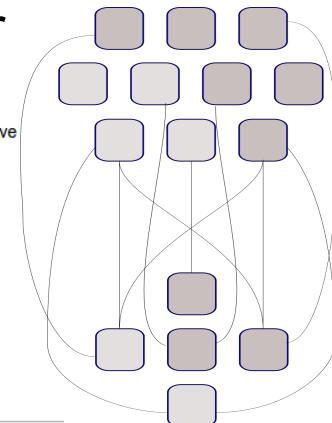


# Dask: Parallel Data Processing

Synthetic views of  
Pandas DataFrames  
with HDFS support



DAG construction and  
workflow manager



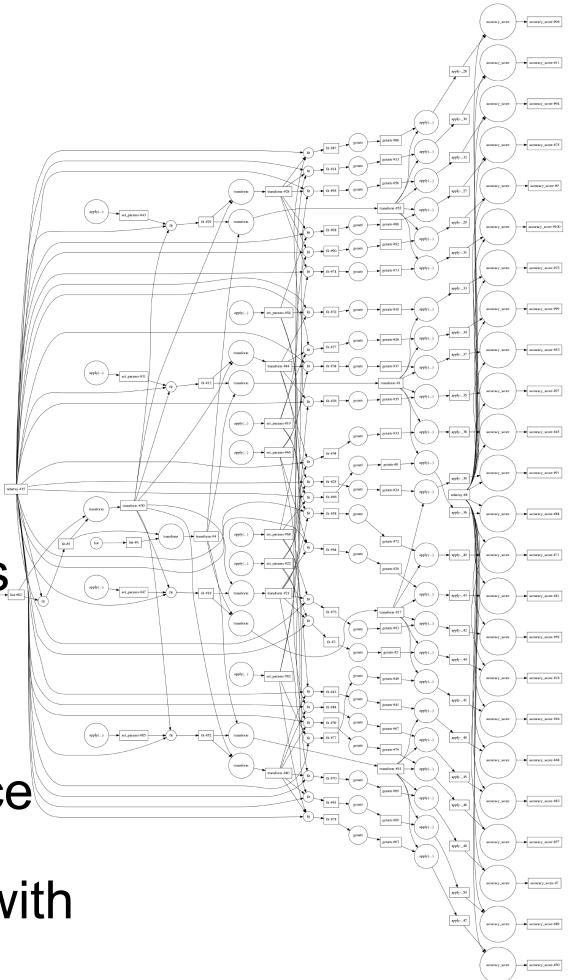
Synthetic views of  
Numpy ndarrays

# Overview of Dask

Dask is a Python parallel computing library that is:

- **Familiar:** Implements parallel NumPy and Pandas objects
- **Fast:** Optimized for demanding for numerical applications
- **Flexible:** for sophisticated and messy algorithms
- **Scales up:** Runs resiliently on clusters of 100s of machines
- **Scales down:** Pragmatic in a single process on a laptop
- **Interactive:** Responsive and fast for interactive data science

Dask **complements** the rest of Anaconda. It was developed with NumPy, Pandas, and scikit-learn developers.



# Dask Collections: Familiar Expressions and API

**Dask array** (mimics NumPy)

```
x.T - x.mean(axis=0)
```

**Dask bag** (collection of data)

```
b.map(json.loads).foldby(...)
```

**Dask dataframe** (mimics Pandas)

```
df.groupby(df.index).value.mean()
```

**Dask delayed** (wraps custom code)

```
def load(filename):  
def clean(data):  
def analyze(result):
```

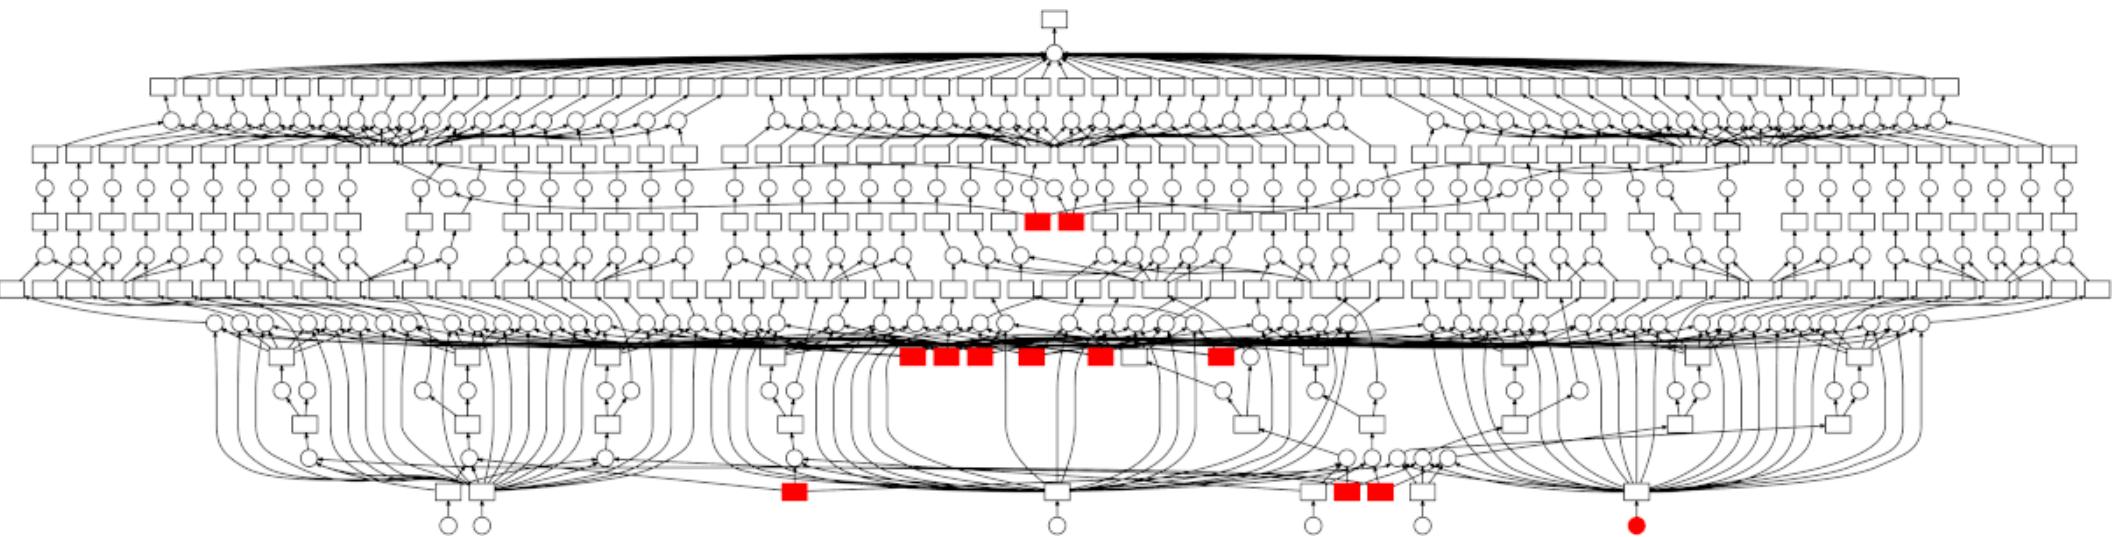
# Dask Dataframes



```
>>> import pandas as pd  
  
>>> df = pd.read_csv('iris.csv')  
  
>>> df.head()  
  
   sepal_length  sepal_width  petal_length  petal_width      species  
0           5.1         3.5          1.4         0.2  Iris-setosa  
1           4.9         3.0          1.4         0.2  Iris-setosa  
2           4.7         3.2          1.3         0.2  Iris-setosa  
3           4.6         3.1          1.5         0.2  Iris-setosa  
4           5.0         3.6          1.4         0.2  Iris-setosa  
  
>>> max_sepallength_setosa = df[df.species  
== 'setosa'].sepal_length.max()  
  
5.799999999999998
```

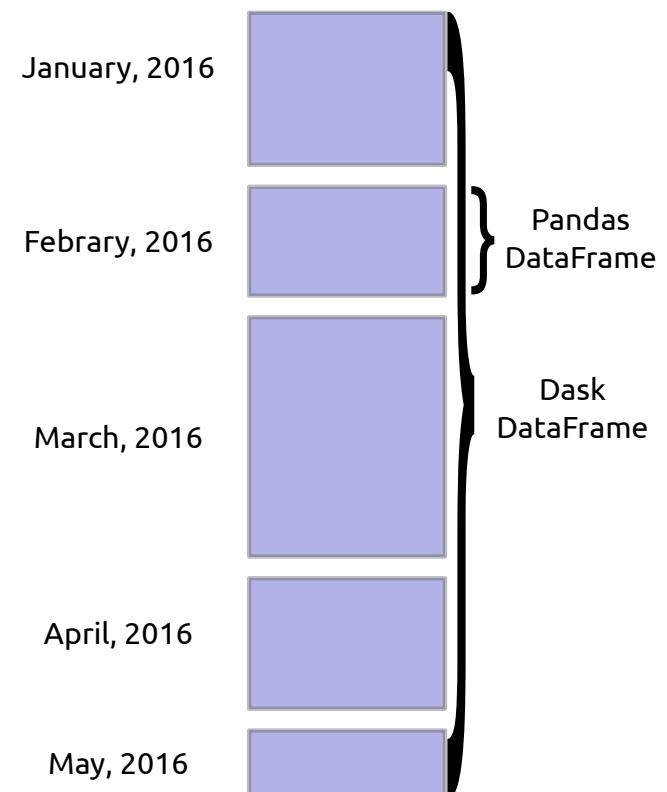
```
>>> import dask.dataframe as dd  
  
>>> ddf = dd.read_csv('*.*csv')  
  
>>> ddf.head()  
  
   sepal_length  sepal_width  petal_length  petal_width      species  
0           5.1         3.5          1.4         0.2  Iris-setosa  
1           4.9         3.0          1.4         0.2  Iris-setosa  
2           4.7         3.2          1.3         0.2  Iris-setosa  
3           4.6         3.1          1.5         0.2  Iris-setosa  
4           5.0         3.6          1.4         0.2  Iris-setosa  
...  
  
>>> d_max_sepallength_setosa = ddf[ddf.species  
== 'setosa'].sepal_length.max()  
  
>>> d_max_sepallength_setosa.compute()  
  
5.799999999999998
```

# Dask Graphs: Example Machine Learning Pipeline



## Example 1: Using Dask DataFrames on a cluster with CSV data

- Built from Pandas DataFrames
- Match Pandas interface
- Access data from HDFS, S3, local, etc.
- Fast, low latency
- Responsive user interface



# Dask Arrays



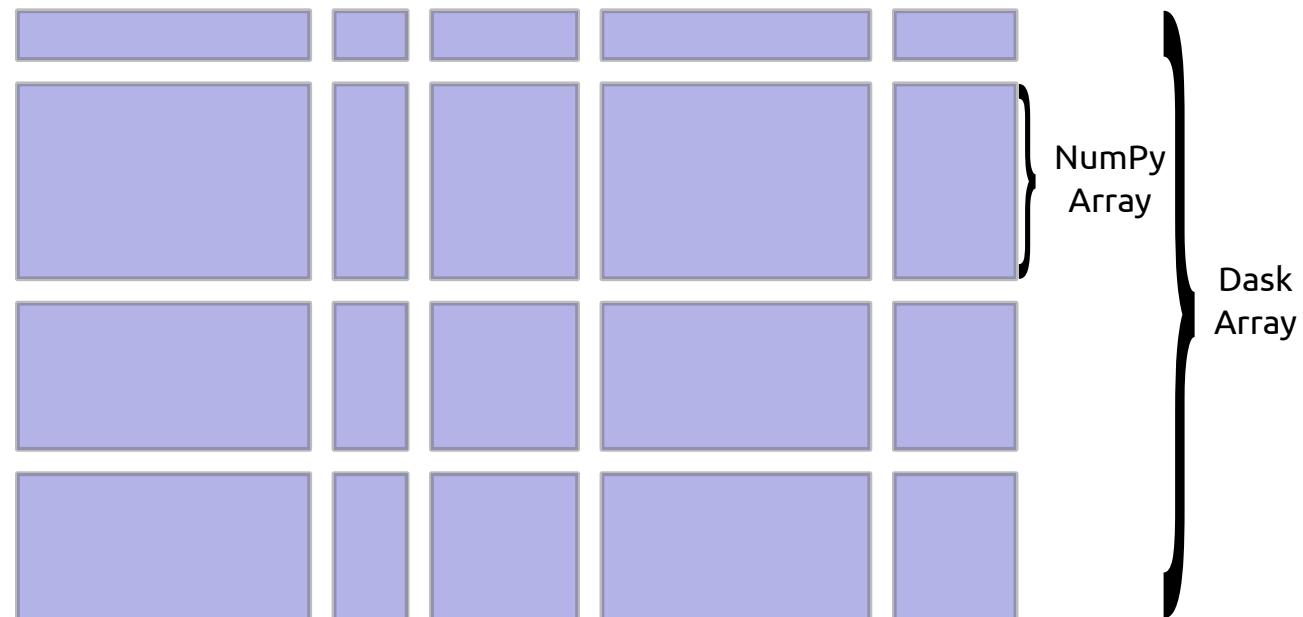
```
>>> import numpy as np
>>> np_ones = np.ones((5000, 1000))
>>> np_ones
array([[ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       ...,
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.]])
>>> np_y = np.log(np_ones + 1)[:5].sum(axis=1)
>>> np_y
array([ 693.14718056,  693.14718056,
       693.14718056,  693.14718056,  693.14718056])
```



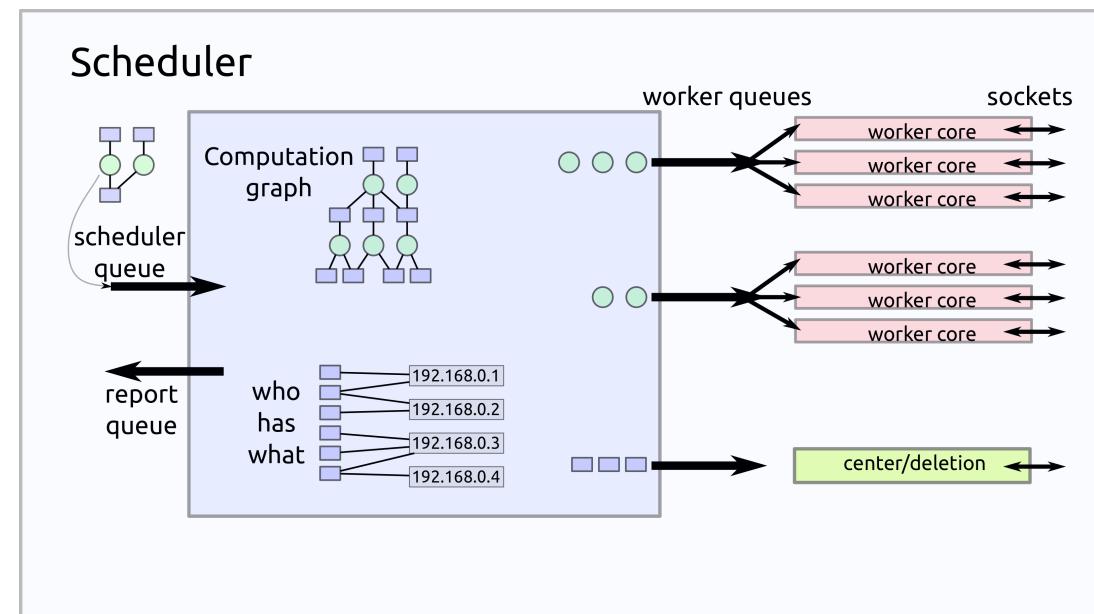
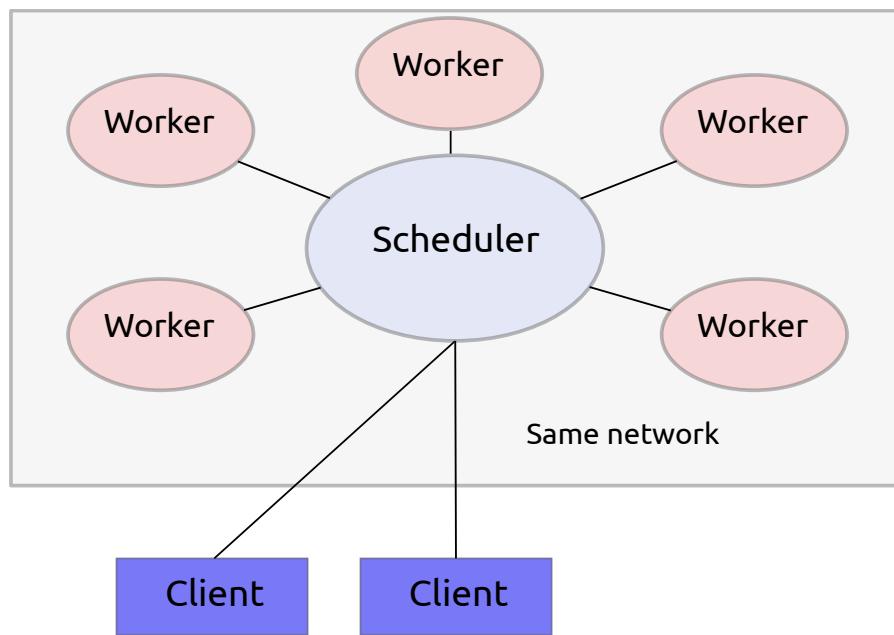
```
>>> import dask.array as da
>>> da_ones = da.ones((5000000, 1000000),
                      chunks=(1000, 1000))
>>> da_ones.compute()
array([[ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       ...,
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.],
       [ 1.,  1.,  1., ...,  1.,  1.,  1.]])
>>> da_y = da.log(da_ones + 1)[:5].sum(axis=1)
>>> np_da_y = np.array(da_y) #fits in memory
array([ 693.14718056,  693.14718056,
       693.14718056,  693.14718056,  693.14718056])
# If result doesn't fit in memory
>>> da_y.to_hdf5('myfile.hdf5', 'result')
```

## Example 3: Using Dask Arrays with global temperature data

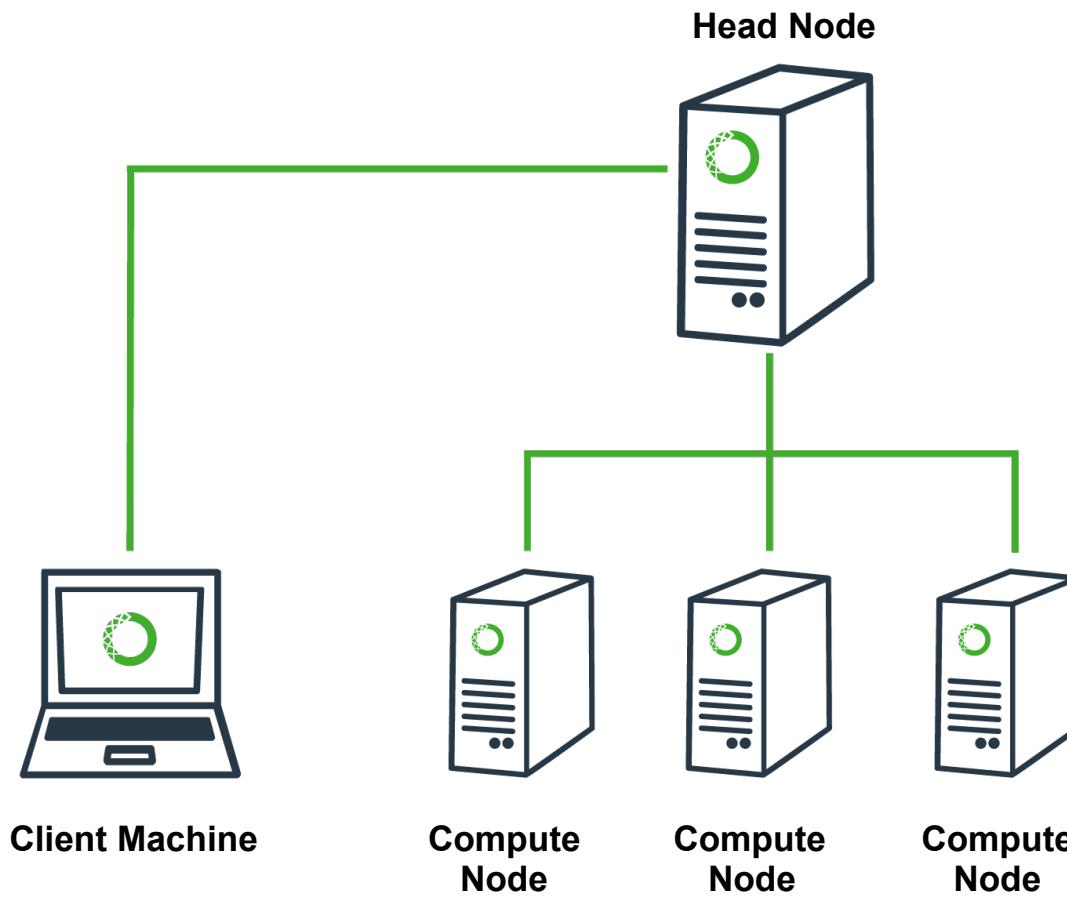
- Built from NumPy n-dimensional arrays
- Matches NumPy interface (subset)
- Solve medium-large problems
- Complex algorithms



# Dask Schedulers: Distributed Scheduler



# Cluster Architecture Diagram

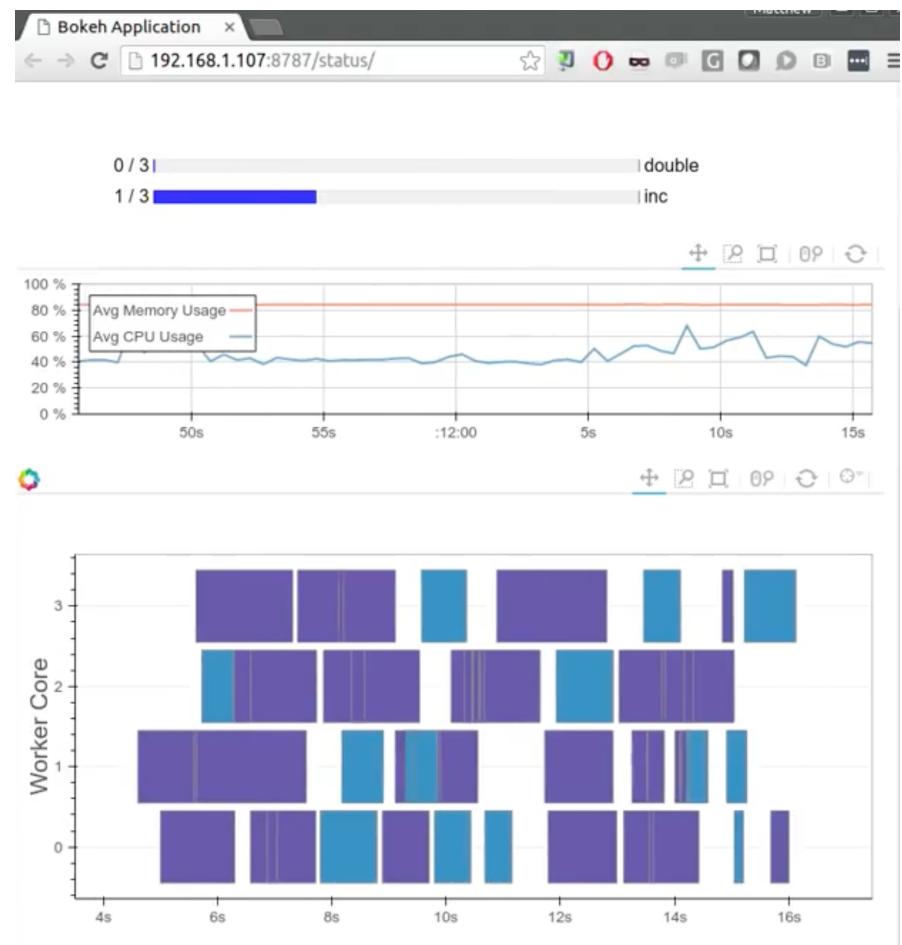
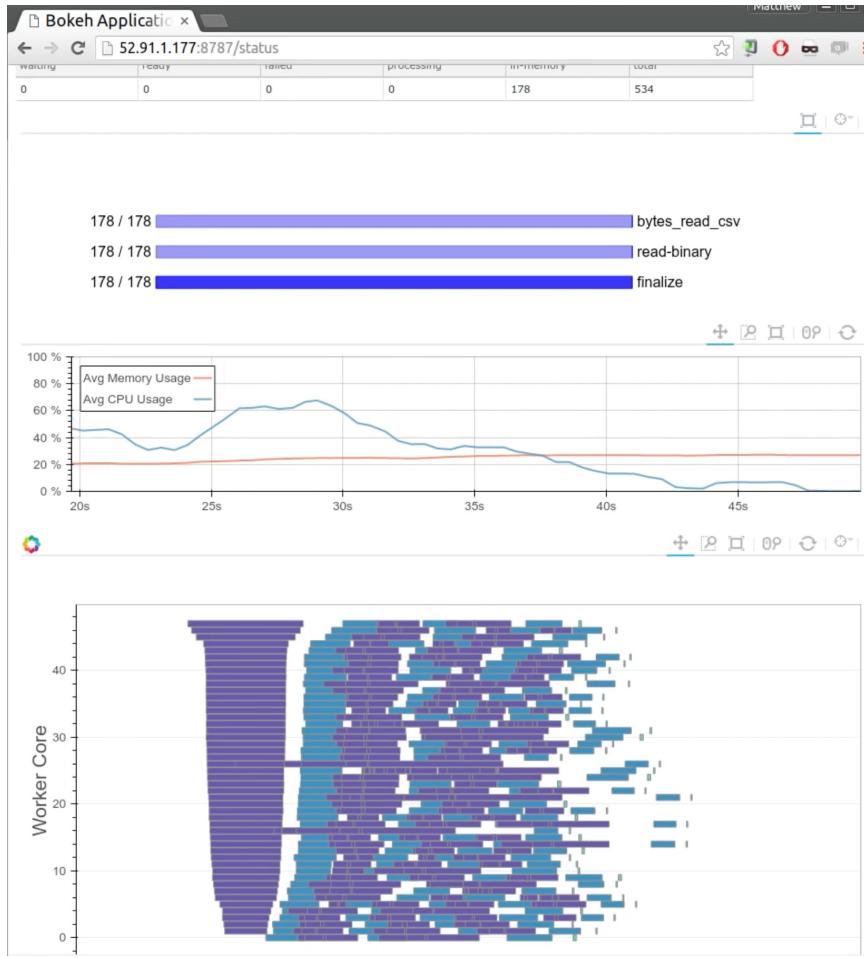


# Using Anaconda and Dask on your Cluster

- Single machine with multiple threads or processes
- On a cluster with **SSH** (dcluster)
- Resource management: **YARN** (knit), **SGE**, **Slurm**
- On the cloud with **Amazon EC2** (dec2)
- On a cluster with **Anaconda for cluster management**
  - Manage multiple conda environments and packages on bare-metal or cloud-based clusters



# Scheduler Visualization with Bokeh



# Numba + Dask

Look at **all** of the data with **Bokeh's datashader**.

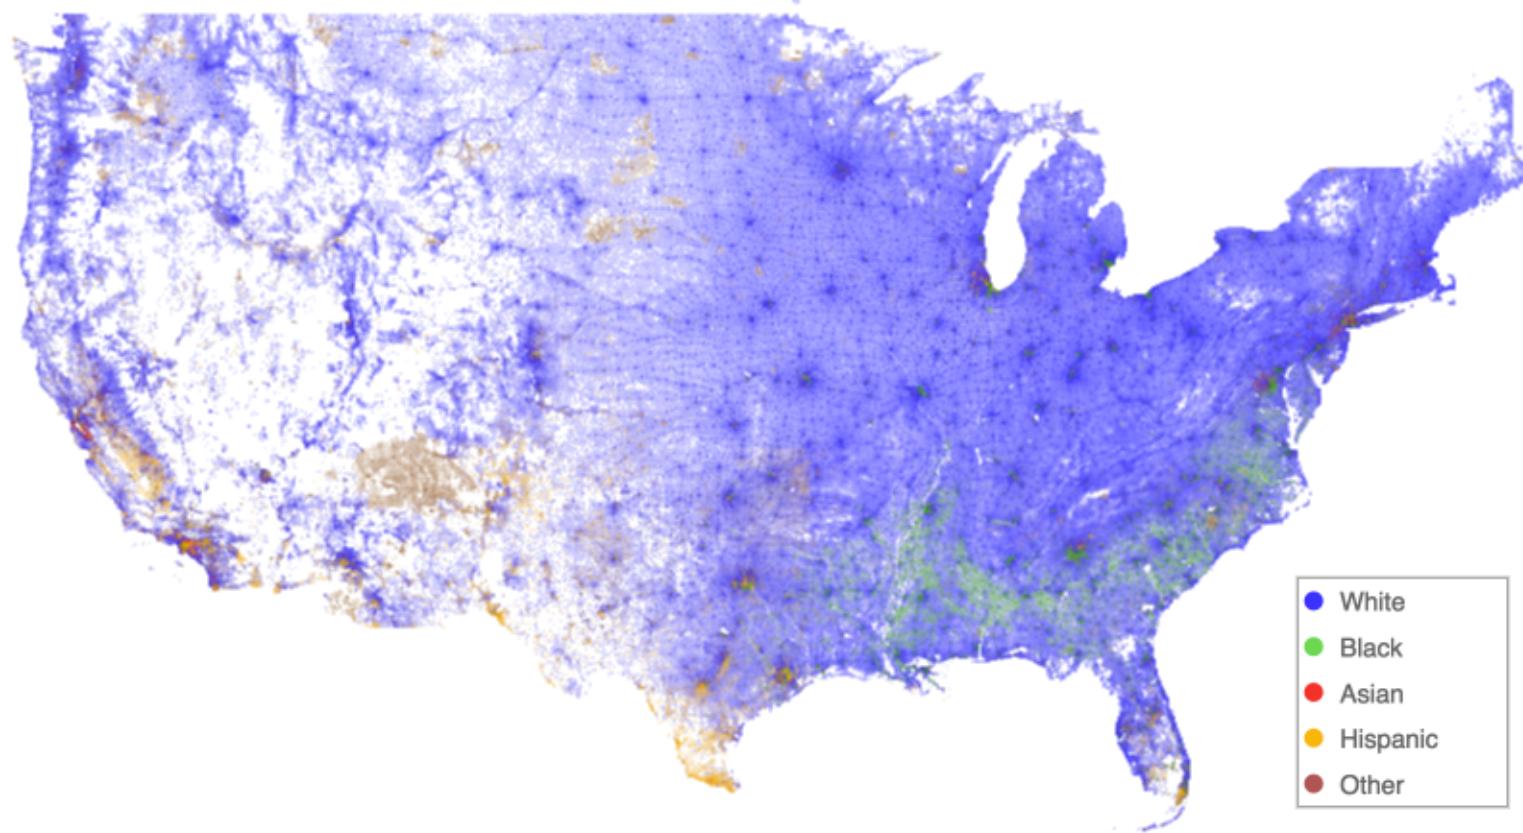
Decouple the data-processing from the visualization. Visualize arbitrarily large data.



e.g. Open Street Map data:

- About **3 billion** GPS coordinates
- <https://blog.openstreetmap.org/2012/04/01/bulk-gps-point-data/>.
- This image was rendered in one minute on a standard MacBook with 16 GB RAM
- Renders in less than a second on several 128GB Amazon EC2 instances

# Categorical data: 2010 US Census



- One point per person
- 300 million total
- Categorized by race
- Interactive rendering with Numba+Dask
- No pre-tiling

# Anaconda and Hadoop

# High Performance Hadoop

## Bottom Line

**2X-100X** faster overall performance

- Interact with data in HDFS and Amazon S3 natively from Python
- Distributed computations without the JVM & Python/Java serialization
- Framework for easy, flexible parallelism using directed acyclic graphs (DAGs)
- Interactive, distributed computing with in-memory persistence/caching

Python & R ecosystem  
MPI

## Bottom Line

- Leverage Python & R with Spark



High Performance,  
Interactive,  
Batch Processing

NumPy, Pandas, ...  
720+ packages

JVM



Batch Processing



PySpark & SparkR



Interactive Processing



Native  
read & write

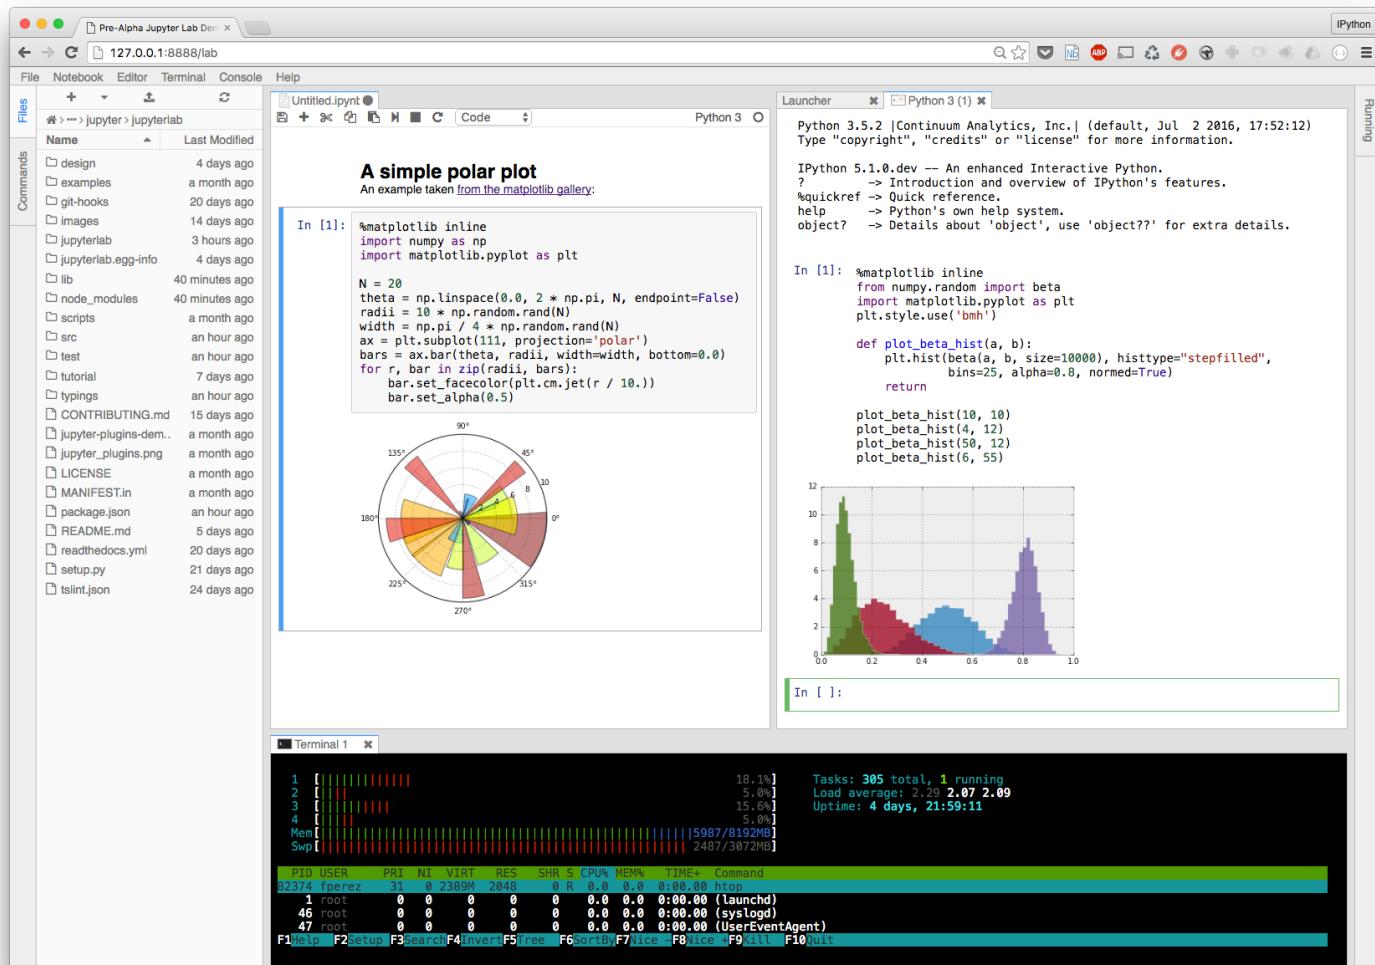
YARN

HDFS

# Jupyterlab (Where Jupyter is heading)

# JupyterLab

- IDE
- Extensible
- Notebook -> Applications



# More Than Just Notebooks

The image displays four screenshots of the Jupyter Notebook interface, illustrating its capabilities beyond just note-taking:

- File Management:** A screenshot of the Jupyter file browser showing the directory structure of a GitHub repository named "altair". The browser includes a search bar, filters for "Files", "Running", and "Clusters", and a toolbar with "Upload", "New", and "Search" buttons.
- Code Execution:** A screenshot of a Jupyter notebook cell showing Python code for reading a setup file and extracting the package version. The code uses regular expressions to find the version number in the file content.
- Terminal Integration:** A screenshot of a terminal window within the Jupyter interface, showing the user navigating through a directory and listing files related to the Altair project.
- Data Visualization:** A screenshot of a Jupyter notebook cell titled "Quick Altair example". It shows a scatter plot generated by Altair, visualizing data from the "cars" dataset. The plot has "Horsepower" on the x-axis and "Miles\_per\_Gallon" on the y-axis. Data points are colored by "Origin" (Europe, Japan, USA).

# Building Blocks

File Browser

Notebooks

Text Editor

Widgets

Output

Terminal

# 2015 User Experience Survey

- Mostly daily/weekly users
- Love the notebook workflow and user experience
- Top needs:
  - Integration with version control systems (git/GitHub)
  - Code/text editing
  - Layout/integration of building blocks
  - Debugger, profiler, variable inspector, etc.

[https://github.com/jupyter/design/blob/master/surveys/2015-notebook-ux/analysis/report\\_dashboard.ipynb](https://github.com/jupyter/design/blob/master/surveys/2015-notebook-ux/analysis/report_dashboard.ipynb)

# A completely modular architecture

The screenshot shows the Jupyter Lab interface with several windows open:

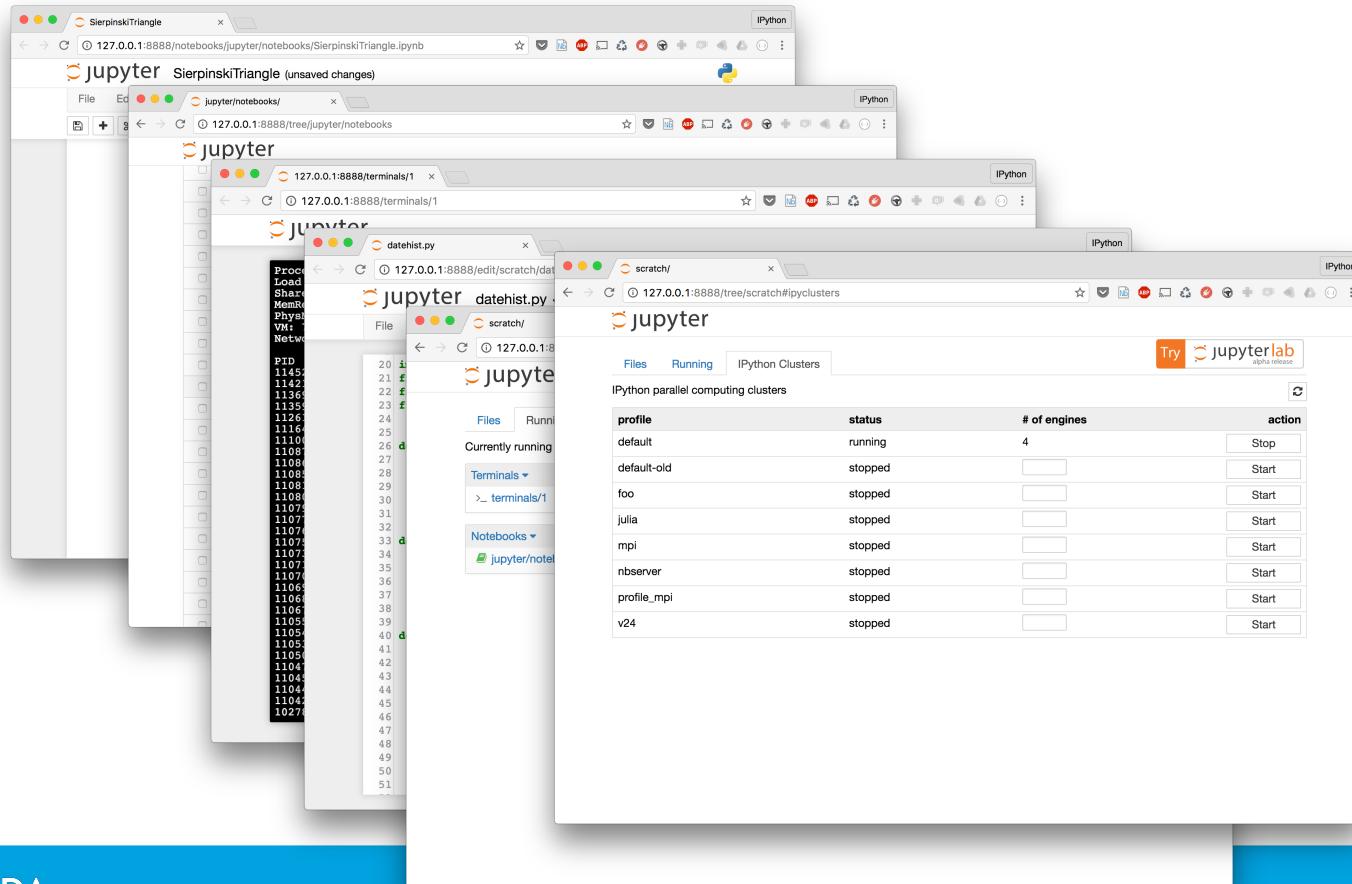
- File**: Shows options like Pre-Alpha Jupyter Lab Dev, 127.0.0.1:8888/lab, File, Notebook, Editor, Terminal, Console, Help.
- Commands**: A sidebar with various keyboard shortcuts for console, editor, file operations, and help.
- CONSOLE**: Displays introductory text about IPython's features and a code cell for plotting beta distributions.
- EDITOR**: Displays a Python script named `mri_with_eeg.py` containing code for loading MRI and EEG data and plotting them.
- FILE OPERATIONS**: Shows commands for closing files, saving documents, and getting help.
- HELP**: Provides links to various Python and Jupyter documentation.
- IMAGE WIDGET**: Shows a histogram of MRI intensity and four EEG signal plots over time.
- Launcher**: Shows the running processes: `mri_with_eeg.py`.

# JupyterLab

- JupyterLab is the natural evolution of the Jupyter Notebook user interface
- JupyterLab is an extensible IDE: *Interactive Data Environment*
- Flexible user interface for assembling the fundamental building blocks of interactive computing
- Modernized JavaScript architecture based on npm/webpack, plugin system, model/view separation
- Built using PhosphorJS (<http://phosphorjs.github.io/>)
- Design-driven development process

<https://github.com/jupyter/jupyterlab>

# The “Notebook”?



# Data Science Workflow

## Exploratory Data Analysis and Machine Learning

